ABSTRACT

Early diagnosis of chronic kidney disease (CKD) is crucial for timely treatment and improved patient outcomes. Chronic kidney disease is a progressive and irreversible condition that affects millions of individuals worldwide and often remains asymptomatic in its early stages, making accurate automated detection a matter of clinical urgency. The World Health Organization (WHO) has reported that CKD currently affects more than 850 million individuals worldwide, with prevalence rates continuing to rise in parallel with the growing global burden of diabetes mellitus, hypertension, and cardiovascular comorbidities. Despite its enormous clinical and economic burden, CKD remains chronically underdiagnosed in its early stages because the insidious nature of the disease means patients often remain asymptomatic until significant kidney function has been lost — frequently up to 60-70% of renal capacity — at which point intervention options are limited and outcomes are substantially worse.

This dissertation investigates the performance of various supervised machine learning models including Logistic Regression, Random Forest, AdaBoost, Gradient Boosting, and XGBoost for CKD classification using a publicly available benchmark dataset comprising 400 patient records with 25 clinical and physiological attributes. The dataset, originally collected from patients in India over a two-month period, captures the diverse spectrum of CKD indicators including hematological markers (hemoglobin, packed cell volume, red cell count, white cell count), biochemical markers (blood urea, serum creatinine, sodium, potassium, blood glucose), urinalysis parameters (specific gravity, albumin, sugar, pus cells, bacteria, red blood cells), physical measurements (age, blood pressure), and comorbidity indicators (hypertension, diabetes mellitus, coronary artery disease, anemia, pedal edema, appetite).

A systematic data preprocessing pipeline was employed to handle missing values, remove high-missingness features, and ensure dataset consistency and completeness. Features with substantial missing data including red blood cells (38% missing), red cell count (32.75% missing), white cell count (26.50% missing), potassium (22% missing), and sodium (21.75% missing) were eliminated to prevent imputation-induced bias. Remaining numerical features with moderate missingness (hemoglobin, blood glucose random, packed cell volume) were imputed using the mean strategy to preserve their distribution, while categorical features were imputed using the mode strategy to retain representative category information.

Principal Component Analysis (PCA) was applied to reduce the dimensionality of the feature space while retaining maximum discriminative information, improving both model interpretability and

computational efficiency. The first two principal components were extracted and visualized to confirm clear class separability between CKD and non-CKD patients. The PCA component loading analysis revealed that hemoglobin (hemo), packed cell volume (pcv), red cell count (rc), and the target classification variable are the most influential features in PC-1, while red blood cells (rbc), pus cells (pc), pus cell clumps (pcc), and bacteria (ba) dominate PC-2. This feature importance insight provides clinically relevant information that aligns with known medical indicators of kidney dysfunction.

The classification models were rigorously evaluated using standard performance metrics including accuracy, precision, recall, and F1-score on an 80:20 train-test split with stratified random sampling to preserve class distributions. The experimental results demonstrate that ensemble-based methods, particularly Random Forest and AdaBoost, achieved the highest accuracy of 97.50% and an F1-score of 0.9750, outperforming both Logistic Regression (96.25%) and Gradient Boosting (93.75%). The near-identical training and testing accuracies for the best-performing models confirm excellent generalization with minimal overfitting, validating the clinical reliability of these approaches.

These findings confirm the robustness and reliability of ensemble learning approaches for medical diagnosis applications and highlight their significant potential for clinical decision support in CKD detection. The correlation matrix analysis further confirms strong inter-feature relationships, particularly between hemoglobin and packed cell volume (correlation coefficient 0.90) and between serum creatinine and blood urea (0.59), which are well-established clinical correlates of renal function. This dissertation presents a comprehensive, reproducible, and clinically interpretable machine learning framework for CKD prediction that can serve as the foundation for an intelligent clinical decision support system (CDSS) to assist healthcare professionals in the early detection and management of chronic kidney disease.

Index Terms: Chronic Kidney Disease (CKD), Machine Learning, Ensemble Learning, Principal Component Analysis (PCA), Random Forest, AdaBoost, Gradient Boosting, Logistic Regression, Clinical Decision Support, Medical Diagnosis.

CHAPTER 1: INTRODUCTION

1.1 Background and Motivation

Chronic kidney disease (CKD) is one of the most significant global health challenges of the twenty-first century. Characterized by a progressive, irreversible loss of kidney function over a period of months or years, CKD silently advances through five stages before culminating in end-stage renal disease

(ESRD), which requires dialysis or transplantation for patient survival. The World Health Organization (WHO) has reported that CKD affects more than 850 million individuals worldwide, with prevalence rates continuing to rise in parallel with the growing burden of diabetes mellitus, hypertension, and cardiovascular comorbidities. These conditions are both primary causes and complications of CKD, creating a complex, bidirectional clinical relationship that further complicates diagnosis and management.

The economic burden of CKD is staggering. In developed countries, the cost of treating end-stage renal disease consumes a disproportionately large fraction of healthcare budgets, with dialysis alone costing tens of thousands of dollars per patient annually. In developing countries, where access to dialysis and transplantation is limited, CKD-related mortality is even higher because affordable late-stage treatment options are unavailable. The Indian context, from which the dataset used in this study originates, illustrates these challenges acutely: with a population exceeding 1.4 billion people and rising rates of diabetes and hypertension, India faces a CKD epidemic that places enormous strain on its healthcare infrastructure and patient families.

Despite its enormous clinical and economic burden, CKD remains chronically underdiagnosed in its early stages. The insidious nature of the disease means that patients often remain asymptomatic until significant kidney function has been lost — frequently up to 60-70% of renal capacity — at which point intervention options are limited and outcomes are substantially worse. Traditional diagnostic pathways rely on biochemical indicators such as serum creatinine, blood urea nitrogen, and glomerular filtration rate (GFR), supplemented by urinalysis findings such as proteinuria and hematuria. While these markers are valuable, they often fail to detect CKD in its earliest stages, particularly in elderly patients, patients with reduced muscle mass, and those with non-diabetic etiologies. Serum creatinine, for example, can remain within the "normal" laboratory reference range until substantial nephron loss has already occurred, providing late warning rather than early detection.

The convergence of large-scale electronic health record (EHR) systems, affordable computational resources, and sophisticated machine learning algorithms creates an unprecedented opportunity to transform CKD diagnostics. Machine learning models can automatically identify complex, non-linear relationships among clinical variables that would be invisible to traditional statistical approaches, potentially enabling earlier and more accurate identification of at-risk patients. Unlike rule-based diagnostic criteria that consider variables independently, machine learning models can simultaneously integrate dozens of features and learn the subtle interaction patterns that characterize early CKD. This dissertation presents a comprehensive study of machine learning-based CKD prediction using a benchmark clinical dataset, with particular emphasis on ensemble techniques and principal component analysis (PCA) for dimensionality reduction and interpretability.

The motivation for this research extends beyond academic interest to direct clinical impact. Every CKD case detected early and managed appropriately represents a patient potentially spared from dialysis dependency, transplantation, or premature death. Even modest improvements in screening sensitivity, when applied at population scale, translate to substantial reductions in disease burden and healthcare costs. Furthermore, accurate ML-based screening tools democratize access to high-quality renal assessment, particularly in resource-limited settings where specialist nephrologists are scarce. A well-trained predictive model deployed on a tablet or smartphone could provide preliminary CKD risk assessment in rural clinics, primary care offices, and community health centers, dramatically expanding the reach of nephrology expertise.

1.2 Chronic Kidney Disease: Clinical Overview

1.2.1 Definition and Staging

Chronic kidney disease is formally defined as an abnormality of kidney structure or function, present for more than three months, with implications for health. The National Kidney Foundation's Kidney Disease Outcomes Quality Initiative (KDOQI) and the more recent KDIGO (Kidney Disease: Improving Global Outcomes) guidelines classify CKD into five stages based on the estimated glomerular filtration rate (eGFR), which serves as the primary measure of kidney function. The stages are:

- **Stage 1**: eGFR ≥ 90 mL/min/1.73m² — Normal or high kidney function with evidence of kidney damage (typically albuminuria or structural abnormalities visible on imaging)
- **Stage 2**: eGFR 60-89 mL/min/1.73m² — Mildly decreased kidney function with persistent kidney damage
- **Stage 3a**: eGFR 45-59 mL/min/1.73m² — Mildly to moderately decreased kidney function
- **Stage 3b**: eGFR 30-44 mL/min/1.73m² — Moderately to severely decreased kidney function
- **Stage 4**: eGFR 15-29 mL/min/1.73m² — Severely decreased kidney function, with patients often requiring nephrology referral and preparation for renal replacement therapy
- **Stage 5**: eGFR < 15 mL/min/1.73m² — Kidney failure (ESRD), typically requiring dialysis or transplantation

Early stages (1-3) are typically asymptomatic, underscoring the critical need for proactive screening and automated detection tools. The dataset used in this study provides binary classification (CKD vs. non-CKD) that corresponds to the clinically critical task of distinguishing patients with any stage of CKD from healthy individuals. While binary classification simplifies the task compared to multi-stage staging, it captures the most important clinical decision point: whether or not a patient requires further nephrological investigation and management.

The progression of CKD is driven by a vicious cycle of nephron loss, glomerular hypertension, and tubular damage. As functional nephrons are progressively lost to disease, the remaining nephrons compensate by hyperfiltration, which paradoxically accelerates their own destruction. This pathophysiological mechanism explains why CKD typically progresses in a non-linear, accelerating fashion once a critical threshold of nephron loss has been crossed, making early detection all the more important.

1.2.2 Etiology and Risk Factors

The primary causes of CKD globally include diabetic nephropathy (accounting for approximately 40% of cases in developed countries), hypertensive nephrosclerosis (approximately 25-30%), glomerulonephritis (10-15%), and polycystic kidney disease and other inherited disorders (5-10%). Secondary risk factors include obesity, smoking, cardiovascular disease, older age, family history of kidney disease, and exposure to nephrotoxic medications including certain antibiotics, contrast agents, and non-steroidal anti-inflammatory drugs (NSAIDs).

The dataset employed in this research captures many of these clinically relevant attributes. The age feature reflects the well-established increase in CKD prevalence with advancing age. Blood pressure (bp) captures the contribution of hypertension. Blood glucose random (bgr) and the categorical diabetes mellitus indicator (dm) reflect diabetic nephropathy risk. Serum creatinine (sc) and blood urea (bu) represent direct biochemical evidence of impaired renal function. Hemoglobin (hemo) reflects the anemia of chronic kidney disease, which results from reduced erythropoietin production by failing kidneys. The categorical comorbidity indicators for hypertension (htn), diabetes mellitus (dm), coronary artery disease (cad), anemia (ane), and pedal edema (pe) capture the multi-system clinical manifestations that frequently accompany CKD.

The presence of multiple comorbidity indicators in the dataset reflects the reality of CKD as a multi-system disease that rarely occurs in isolation. The strong association between CKD and cardiovascular disease, in particular, reflects shared risk factors (hypertension, diabetes, dyslipidemia) and the bidirectional pathophysiological relationship in which each condition accelerates the other. Patients with CKD have a 5-10 fold increased risk of cardiovascular mortality compared to age-matched controls, often dying from heart disease before progressing to end-stage renal failure.

1.2.3 Clinical Indicators and Biomarkers

The biochemical and hematological indicators captured in the CKD dataset reflect the multisystem impact of kidney dysfunction. Elevated serum creatinine and blood urea reflect impaired nitrogenous waste excretion. Reduced hemoglobin and packed cell volume indicate the anemia of chronic kidney disease, caused by decreased erythropoietin production by failing kidneys and shortened red blood

cell survival. Abnormal specific gravity, albumin, and sugar levels in urine reflect disrupted tubular reabsorption and glomerular filtration.

The specific gravity of urine (sg), albumin (al), and urine sugar (su) are key urinalysis findings. Low specific gravity (close to 1.010, indicating isosthenuria) suggests loss of urinary concentrating ability, an early sign of tubular dysfunction. Albuminuria reflects glomerular injury and is a strong predictor of CKD progression. Glycosuria suggests either uncontrolled diabetes mellitus or proximal tubular dysfunction (renal glycosuria). Blood pressure (bp) reflects hypertensive nephropathy, which both causes and results from CKD.

The presence of pus cells (pc), pus cell clumps (pcc), bacteria (ba), and abnormal red blood cells (rbc) in urine indicate infectious or inflammatory kidney pathology. Pyuria with bacteriuria suggests urinary tract infection, which can be both a cause and complication of CKD. Pus cell clumps may indicate pyelonephritis (kidney infection). Hematuria with abnormal red blood cell morphology (dysmorphic red cells, red cell casts) suggests glomerulonephritis. The combination of these urinalysis findings with the biochemical and hematological markers provides a comprehensive clinical picture that the machine learning model must learn to interpret.

The serum electrolyte abnormalities captured in the dataset (sodium and potassium, though these were excluded due to high missingness) reflect the kidneys' role in fluid and electrolyte homeostasis. Hyponatremia in CKD typically reflects volume overload, while hyperkalemia is a dangerous complication of advanced CKD that can cause cardiac arrhythmias. The exclusion of these features due to missing data represents a limitation of the current study that should be addressed in future work with more complete datasets.

1.3 Role of Machine Learning in Medical Diagnosis

Machine learning has emerged as a transformative technology in medical diagnostics, offering the ability to process large, complex, multi-dimensional datasets and extract predictive patterns that exceed the capability of traditional clinical decision rules. In the domain of CKD, ML algorithms can integrate diverse biochemical, hematological, and physiological parameters into a unified predictive model, potentially enabling earlier diagnosis with greater sensitivity and specificity than conventional threshold-based clinical criteria.

Several categories of ML algorithms have demonstrated effectiveness for CKD prediction. Linear models such as Logistic Regression offer interpretability and clinical transparency, with model coefficients directly indicating how each feature contributes to the predicted probability of disease. Tree-based ensemble methods including Random Forest, Gradient Boosting, AdaBoost, and XGBoost combine the predictions of multiple decision trees to achieve superior accuracy and robustness, while

still providing some interpretability through feature importance scores and partial dependence plots. Deep learning approaches, while powerful, typically require larger datasets and offer reduced interpretability, making them less immediately practical for clinical deployment compared to well-validated ensemble methods.

The integration of dimensionality reduction techniques such as Principal Component Analysis (PCA) with ML classifiers further enhances model performance by reducing feature redundancy, mitigating the curse of dimensionality, and improving computational efficiency. PCA transforms the original correlated feature space into uncorrelated principal components ordered by explained variance, enabling more stable and generalizable model training. The ability to visualize high-dimensional data in two or three principal component dimensions also facilitates exploratory data analysis and quality control.

The fundamental advantage of machine learning over traditional statistical approaches lies in the ability to capture non-linear relationships and high-order feature interactions automatically, without requiring the analyst to specify these relationships in advance. For CKD prediction, this is particularly valuable because the relationship between individual clinical features and disease status is rarely linear or monotonic. For example, both extremely high and extremely low serum creatinine values may indicate pathology, while values within a wide intermediate range are normal. Traditional logistic regression cannot capture this U-shaped relationship without explicit transformation, but tree-based ensemble methods learn it automatically from the data.

Furthermore, machine learning excels at handling the heterogeneity of clinical populations. Different patients with CKD present with different combinations of clinical features depending on the underlying etiology (diabetic vs. hypertensive vs. glomerulonephritic), the stage of disease, the presence of comorbidities, and individual genetic factors. A well-trained ML model implicitly learns to recognize multiple distinct patterns of CKD presentation and assigns appropriate diagnostic probability to each, whereas rule-based systems typically work best for the "average" case and perform poorly on atypical presentations.

1.4 Research Insights

This research is motivated by the critical need for accurate, interpretable, and computationally efficient CKD prediction models that can be deployed in clinical settings with limited computational resources. The benchmark CKD dataset from the UCI Machine Learning Repository, collected from patients in India over a two-month period, provides a rich set of 25 clinical and physiological attributes that capture the diverse spectrum of CKD indicators. Despite the widespread use of this dataset in the literature, a systematic comparative assessment of ensemble and non-ensemble techniques

incorporating PCA-based dimensionality reduction with consistent preprocessing has remained insufficient.

This study addresses this gap by presenting a unified analytical framework that evaluates four prominent classifiers — Logistic Regression, Random Forest, AdaBoost, and Gradient Boosting — on this benchmark dataset, offering a comprehensive evaluation of their effectiveness for CKD diagnosis. The use of PCA provides both dimensionality reduction benefits and enhanced feature interpretability through component loading analysis, revealing which clinical variables are most discriminative for CKD classification. By using a consistent preprocessing pipeline, identical train-test splits, and uniform evaluation metrics across all classifiers, this study enables a fair and rigorous comparison that has been lacking in much of the prior literature.

The clinical insights generated by this research extend beyond pure model performance. The PCA component loading analysis identifies which clinical features carry the most discriminative information for CKD detection, providing valuable guidance for clinicians regarding which laboratory tests are most important to order and interpret. The correlation analysis reveals the underlying structure of clinical relationships in the dataset, confirming established medical knowledge (such as the strong correlation between hemoglobin and packed cell volume) while potentially highlighting unexpected associations that warrant further investigation.

The comparative analysis of ensemble vs. non-ensemble methods provides practical guidance for healthcare informatics practitioners selecting algorithms for clinical deployment. The finding that Random Forest and AdaBoost achieve identical performance (97.50%) despite using fundamentally different ensemble mechanisms (bagging vs. boosting) suggests that both algorithms have effectively saturated the discriminative information available in this dataset, and that further accuracy improvements may require either richer feature sets or more sophisticated algorithms such as deep learning.

1.5 Problems Identified

Several specific problems in the current state of CKD machine learning research are addressed by this dissertation:

- Traditional diagnostic methods for CKD rely on single-threshold biochemical criteria such as serum creatinine cutoffs that often fail to detect disease in asymptomatic early stages, particularly in elderly patients and those with reduced muscle mass where creatinine is a less reliable marker.
- High-dimensional clinical datasets with missing values, mixed data types (numerical and categorical), and class imbalance present significant challenges for conventional statistical classifiers that assume complete data and homogeneous feature types.

- Many existing ML approaches for CKD lack systematic preprocessing pipelines, leading to biased or unreliable model training when missing values are handled inappropriately or categorical features are encoded inconsistently.
- Gradient Boosting and other complex ensemble methods show sensitivity to hyperparameter settings, potentially producing suboptimal performance without careful tuning, but published studies often report results from default hyperparameters without systematic optimization.
- The interpretability of black-box ML models remains a significant barrier to clinical adoption, particularly in regulated medical environments where explainability is a regulatory requirement.
- Minor class imbalance in clinical datasets can lead to undefined metrics and biased recall evaluation, particularly for rare disease categories or extreme test set splits.
- Many published CKD ML studies use inconsistent evaluation methodologies (different splits, different metrics, different preprocessing) that make direct comparison of results difficult or impossible.

1.6 Research Gaps

The literature review (presented in Chapter 2) identified several specific research gaps that this dissertation addresses:

- Existing studies rarely provide a systematic, unified comparison of both ensemble and non-ensemble methods with consistent preprocessing and identical evaluation conditions on the same benchmark dataset.
- The combined effect of PCA-based dimensionality reduction with ensemble classifiers on CKD prediction performance has not been comprehensively explored, with most studies applying either PCA or ensemble methods, but not both in a coordinated framework.
- Most published work does not provide detailed component loading analysis to identify the clinical interpretability of PCA dimensions in CKD datasets, missing an opportunity to bridge the gap between data-driven dimensionality reduction and clinical knowledge.
- Hyperparameter optimization for Gradient Boosting and XGBoost classifiers specifically for the CKD benchmark dataset remains underexplored, leaving open the question of whether these methods can match or exceed Random Forest performance with proper tuning.
- Integration of explainable AI techniques with ensemble CKD classifiers for clinical trust and regulatory compliance has not been widely implemented or systematically evaluated.
- The reproducibility of published CKD ML results is often limited by incomplete documentation of preprocessing steps, hyperparameter settings, and evaluation procedures, making it difficult for subsequent researchers to build upon prior work.

1.7 Objectives

The primary objectives of this research are:

1. To construct machine learning models that accurately assess CKD diagnosis using a benchmark clinical dataset of 400 patient records with 25 physiological attributes.
2. To design and implement a systematic data preprocessing pipeline including missing value imputation, categorical encoding, normalization, and feature elimination for high-missingness variables, ensuring data integrity and reproducibility.
3. To apply Principal Component Analysis (PCA) for dimensionality reduction and feature interpretability, extracting the two most discriminative principal components from the clinical dataset and analyzing their component loadings for clinical relevance.
4. To evaluate and compare ensemble and non-ensemble machine learning approaches including Logistic Regression, Random Forest, AdaBoost, and Gradient Boosting using accuracy, precision, recall, and F1-score under identical experimental conditions.
5. To identify the most effective machine learning approach for CKD detection and demonstrate the contribution of PCA to improved computational effectiveness and model interpretability.
6. To provide a foundation for a reliable clinical decision support system (CDSS) that can assist healthcare professionals in early CKD recognition and treatment planning, with particular emphasis on deployability in resource-limited settings.

1.8 Organization of the Dissertation

The remainder of this dissertation is organized as follows. Chapter 2 presents a comprehensive review of related literature on machine learning approaches for CKD detection, ensemble methods, and dimensionality reduction techniques, identifying the specific gaps that motivate this work. Chapter 3 describes the dataset used in this study, provides a detailed feature description, and presents exploratory data analysis including visualization of feature distributions and correlation patterns. Chapter 4 details the proposed methodology including the preprocessing pipeline, PCA implementation, classifier descriptions with mathematical foundations, and evaluation metrics. Chapter 5 presents the experimental results and provides a detailed discussion of model performance, PCA component interpretability, comparative analysis, and clinical implications. Chapter 6 concludes the dissertation and outlines directions for future research including hyperparameter optimization, deep learning integration, explainable AI, multi-stage CKD classification, longitudinal monitoring, federated learning, and clinical workflow integration.

CHAPTER 2: LITERATURE REVIEW

2.1 Introduction

The application of machine learning techniques to chronic kidney disease detection and prediction has been an active area of research over the past decade, driven by the clinical urgency of early CKD diagnosis and the increasing availability of large clinical datasets. This chapter surveys the most significant contributions to this field, organized thematically by methodological approach, and identifies the specific research gaps that this dissertation aims to address. The literature review is organized into six main categories: (1) machine learning approaches for CKD detection generally, (2) ensemble-based methods including bagging and boosting, (3) dimensionality reduction in healthcare data, (4) explainable AI in medical classification, (5) hybrid deep learning and machine learning approaches, and (6) hyperparameter optimization and feature selection.

2.2 Machine Learning Approaches for CKD Detection

The earliest applications of machine learning to CKD prediction utilized classical algorithms including Logistic Regression, k-Nearest Neighbors, Naive Bayes, Decision Trees, and Support Vector Machines on small clinical datasets. These foundational studies established the feasibility of ML-based CKD prediction and identified the most discriminative clinical features.

Using ensemble models such as Random Forest and Gradient Boosting, Dahiya et al. [1] demonstrated the reliability of ensemble approaches for medical datasets and achieved an accuracy of approximately 96.8%, validating the strength of tree-based ensemble methods on clinical tabular data. The authors particularly highlighted the robustness of Random Forest against noise and missing values in medical records, attributes that are essential for real-world clinical deployment where data quality cannot be guaranteed. Their work established a baseline performance benchmark that subsequent studies have aimed to match or exceed.

For early-stage CKD recognition, Sharma and Kumar [2] applied Logistic Regression and Support Vector Machines, emphasizing the value of feature selection and optimal preprocessing. Their work demonstrated that simpler linear models, when properly preprocessed, can achieve competitive results with more complex approaches, and underscored the importance of systematic data cleaning in clinical ML applications. They also noted that Logistic Regression's interpretability advantage may make it preferable to more complex methods in clinical settings where regulatory and explainability requirements are paramount, even at the cost of marginal accuracy reduction.

Yadav et al. [3] presented a comparative study of various classifiers including Naive Bayes, k-NN, and Decision Trees, discovering that Random Forest was the most reliable for CKD classification, achieving the highest accuracy among tested models. Their study provided early evidence for the superiority of ensemble methods over individual classifiers for this clinical task. They also analyzed the computational efficiency of different methods, concluding that Random Forest offers the best balance of accuracy and computational cost for typical clinical deployment scenarios.

Khan et al. [12] focused specifically on early-stage CKD detection, recognizing that the clinical value of ML models is highest when they can detect disease before symptoms manifest. Their work used clinical data analysis to identify the earliest detectable patterns of CKD progression and demonstrated that machine learning can outperform traditional eGFR-based screening for early-stage disease identification. This finding has significant clinical implications because traditional creatinine-based screening tends to under-detect early CKD, particularly in populations where serum creatinine is a less reliable marker.

Raj and Menon [13] conducted a comprehensive performance evaluation of multiple ML models for CKD classification, providing systematic benchmarking that subsequent researchers have used as reference. Their evaluation framework, which included multiple datasets and rigorous cross-validation, established methodological standards for the field. They concluded that ensemble methods consistently outperform single classifiers across diverse CKD datasets, validating the generality of this finding beyond any single benchmark.

2.3 Ensemble-Based Methods in Medical Diagnosis

Ensemble methods, which combine the predictions of multiple base learners to achieve superior performance compared to individual classifiers, have emerged as the dominant paradigm for clinical prediction tasks. The two main ensemble paradigms — bagging (bootstrap aggregating) exemplified by Random Forest, and boosting exemplified by AdaBoost, Gradient Boosting, and XGBoost — have each received substantial attention in the CKD literature.

Ahmed et al. [4] proposed a hybrid ensemble model combining AdaBoost and XGBoost algorithms, demonstrating that combining multiple boosting-based learners further improves CKD classification performance compared to either algorithm alone. Their hybrid approach achieved 97.1% accuracy, exceeding the individual performance of either constituent algorithm. The success of this hybrid motivates the comparative investigation of AdaBoost and Gradient Boosting conducted in this dissertation, while suggesting that future work might explore stacking or blending these methods together.

Gupta and Tiwari [7] examined bagging and boosting techniques specifically for CKD, finding that AdaBoost outperforms other boosting methods when handling class imbalance in CKD data. Their analysis provided detailed insights into why AdaBoost is particularly well-suited to clinical datasets: the iterative reweighting mechanism naturally focuses learning capacity on hard-to-classify cases, which often correspond to clinically ambiguous CKD presentations near the diagnostic boundary. This finding is directly consistent with the strong AdaBoost performance observed in the present dissertation.

Roy and Das [10] enhanced chronic kidney disease prediction through an optimized Random Forest model with hyperparameter tuning. Their study demonstrated that careful selection of the number of trees, maximum depth, minimum samples per leaf, and feature subset size at each split can boost Random Forest performance by 1-2 percentage points compared to default settings. Although the present dissertation uses default hyperparameters for fair comparison, future work should incorporate the systematic hyperparameter optimization approaches described by Roy and Das.

Chen and Zhao [11] investigated data-driven CKD detection using an optimized Gradient Boosting framework. Their work specifically addressed the hyperparameter sensitivity of Gradient Boosting that has been noted in multiple studies, demonstrating that with proper learning rate scheduling, tree depth selection, and regularization, Gradient Boosting can match or exceed the performance of Random Forest. Their methodology provides a roadmap for future hyperparameter optimization of the Gradient Boosting baseline in the present study, which underperformed Random Forest at default settings.

Zhang et al. [14] specifically applied XGBoost to CKD detection from medical datasets, leveraging XGBoost's regularization and parallel processing advantages. They achieved 97.4% accuracy and provided detailed feature importance analysis that aligned closely with clinical knowledge. Their study established XGBoost as a competitive alternative to Random Forest for tabular clinical data, particularly when computational efficiency is a concern.

Patel et al. [16] conducted a comparative study of ML and DL models for CKD detection, finding that for the typical clinical dataset size of a few hundred to a few thousand patients, classical ML methods (particularly ensemble methods) consistently match or exceed deep learning performance while requiring substantially less computational resources and providing better interpretability. This finding validates the focus of the present dissertation on classical ensemble methods rather than deep learning architectures.

2.4 Dimensionality Reduction in Healthcare Data

The challenge of high-dimensional clinical data with correlated features and missing values has motivated the application of dimensionality reduction techniques. PCA has been widely adopted as a preprocessing step for medical ML tasks. Zhang et al. [14] emphasized the advantages of integrating

PCA with explainable ensemble learning to improve model interpretability, demonstrating that PCA-reduced features maintain discriminative power while significantly reducing computational complexity. Their work showed that the first few principal components typically capture 80-95% of the variance in clinical datasets while reducing dimensionality by an order of magnitude.

Thomas and Joseph [17] further highlighted the benefits of combining PCA with ensemble classifiers for CKD classification, showing improved generalization and enhanced transparency in model predictions. Their work directly motivates the PCA-ensemble integration approach adopted in this dissertation, particularly the component loading analysis for clinical feature interpretation. They demonstrated that PCA component loadings can be visualized as a heatmap to provide intuitive understanding of which original features drive each principal component, bridging the gap between mathematical dimensionality reduction and clinical interpretability.

Liu et al. [18] developed automated CKD risk assessment using ML-based diagnostic models with dimensionality reduction. Their study demonstrated that beyond performance improvements, dimensionality reduction provides essential benefits for clinical deployment: smaller models load faster, require less memory, and can be deployed on resource-constrained devices such as point-of-care testing equipment in primary care clinics. These deployment advantages are particularly important for the goal of expanding CKD screening to underserved populations and resource-limited settings.

Beyond PCA, alternative dimensionality reduction techniques have been explored in the CKD literature. t-SNE (t-distributed Stochastic Neighbor Embedding) and UMAP (Uniform Manifold Approximation and Projection) provide non-linear dimensionality reduction that can capture more complex patterns than PCA, but at the cost of reduced interpretability and the inability to apply the learned transformation to new data. Autoencoder-based dimensionality reduction has also been investigated, particularly for very high-dimensional datasets, but its effectiveness on small clinical datasets is often limited by overfitting.

Feature selection methods (selecting a subset of original features) provide an alternative to feature extraction methods (creating new features from combinations of originals). Recursive feature elimination, mutual information-based selection, and L1-regularized models such as Lasso regression have all been applied to CKD datasets. These methods preserve direct interpretability (since selected features remain in their original clinical units) but may miss complex feature interactions that PCA can capture in its principal components.

2.5 Explainable AI in CKD Classification

For clinical deployment, model transparency and interpretability are essential. Future implementations should incorporate explainable AI frameworks to provide patient-level and population-level explanations of model predictions. Rahaman et al. [5] introduced explainable AI techniques, specifically SHAP (SHapley Additive exPlanations) and LIME (Local Interpretable Model-agnostic Explanations), for evaluating ensemble model predictions in CKD diagnosis. Their work highlighted the critical importance of model transparency and interpretability for clinical trust and regulatory compliance, demonstrating that even high-performing black-box models require interpretability frameworks before clinical deployment.

The clinical value of XAI techniques extends beyond regulatory compliance to direct patient care. When a clinician can understand why a model predicts CKD for a specific patient — for example, "this patient's elevated serum creatinine and reduced hemoglobin contribute most to the CKD prediction" — they can validate the prediction against their clinical judgment, identify potential model errors, and use the model's reasoning to guide further diagnostic workup. This level of explanation is impossible with raw probability scores alone.

Rahman and Ferdous [19] showed that hyperparameter tuning of decision tree ensembles significantly improves diagnostic performance and model stability. Their systematic study across multiple CKD datasets demonstrated consistent gains from hyperparameter optimization, motivating the inclusion of hyperparameter tuning in the future work proposed in this dissertation. They also provided interpretability analysis showing that optimized models tend to focus on a smaller, more clinically meaningful subset of features compared to default configurations.

Singh and Nair [23] addressed clinical usability and reliability by building an interpretable hybrid ensemble model for CKD classification, demonstrating that interpretability and high accuracy are not mutually exclusive goals. Their hybrid approach combines a high-accuracy black-box ensemble for prediction with a separate interpretable surrogate model for explanation, providing both strong performance and clinical transparency. This approach represents a promising direction for clinical deployment of CKD ML models.

The integration of XAI with the PCA-based approach used in this dissertation is particularly promising because the PCA component loadings already provide a form of explanation: they identify which original features drive each principal component used by the classifier. Future work integrating SHAP or LIME with the PCA-ensemble framework could provide comprehensive explanations at multiple levels of abstraction, from the original features through the principal components to the final prediction.

2.6 Hybrid Deep Learning and Machine Learning Approaches

While classical machine learning has dominated the CKD prediction literature, deep learning approaches have also been explored, particularly in combination with classical methods in hybrid architectures. Li et al. [6] developed a deep learning-assisted clinical decision support system combining Gradient Boosting and Convolutional Neural Networks (CNN), achieving 97.2% accuracy on CKD datasets. Their hybrid approach uses CNN for feature extraction from any image data (such as renal ultrasound or biopsy histology), combined with Gradient Boosting for final prediction integration with tabular clinical data. While powerful, this approach requires substantial computational resources and labeled image data, making it less broadly applicable than tabular-only methods.

Mehta and Patel [8] applied deep neural networks to CKD prediction using only tabular data and achieved 94.8% accuracy. Their study illustrated both the potential and limitations of deep learning for tabular clinical data: while deep neural networks can in principle capture complex patterns, they typically require larger datasets than are available in most CKD studies and offer reduced interpretability compared to ensemble methods. Their results suggest that for the typical CKD dataset size (a few hundred patients), classical ensemble methods remain the optimal choice.

Hossain et al. [15] proposed a hybrid CNN-ML approach for CKD prediction using tabular data. Their architecture uses CNN-style 1D convolutions to learn feature interactions, followed by classical ML classifiers for final prediction. This hybrid approach can capture some non-linear patterns that pure ensemble methods might miss, while still benefiting from the regularization and interpretability of classical methods. They achieved competitive accuracy and provided a template for future hybrid architectures.

Singh et al. [9] explored feature selection-based CKD classification using hybrid ensemble methods. Their work emphasized that hybrid approaches benefit most from careful feature engineering and selection, and that throwing more model complexity at a small clinical dataset is rarely a winning strategy. They achieved strong results with relatively simple feature selection followed by ensemble classification, supporting the methodology adopted in this dissertation.

2.7 Recent Advances and Emerging Approaches

Recent literature has begun to explore more advanced approaches that go beyond classical supervised learning. Nguyen and Tran [20] developed an optimized AdaBoost variant specifically for CKD classification, achieving improved performance through adaptive learning rate scheduling and

weighted sample selection. Their methodological innovations could be incorporated into the AdaBoost baseline used in the present dissertation to potentially exceed the current 97.50% accuracy.

Chen [21] analyzed clinical CKD datasets using improved Gradient Boosting techniques, addressing the hyperparameter sensitivity issue noted in multiple previous studies. Their work provided practical guidance for tuning Gradient Boosting on small clinical datasets, including recommendations for cross-validation strategies, regularization parameter selection, and learning rate scheduling.

Federated learning has emerged as a promising approach for training CKD models across multiple institutions without sharing raw patient data, addressing the privacy and regulatory constraints that limit traditional centralized ML approaches. While not yet widely adopted in CKD research, federated learning is expected to grow in importance as healthcare AI moves toward multi-institutional deployment.

Self-supervised learning and transfer learning approaches are also beginning to appear in the CKD literature, though their application to small tabular clinical datasets remains limited. These approaches show greater promise for image-based medical AI tasks where pre-trained representations from large natural image datasets can transfer to medical imaging tasks.

2.8 Summary of Related Works

The literature reviewed above can be summarized in terms of methodology and reported performance. The accuracy range reported across studies (94-97.5%) reflects both the inherent discriminability of the CKD benchmark dataset and the impact of methodological choices. Studies achieving higher accuracy typically combine careful preprocessing, dimensionality reduction or feature selection, ensemble methods, and hyperparameter optimization. The present dissertation incorporates the first three elements and identifies hyperparameter optimization as a future direction.

Table 2.1 Summary of Related Works in CKD Classification

Reference	Method	Accuracy	Key Contribution
Dahiya et al. [1]	RF + Gradient Boosting	96.8%	Ensemble reliability for medical data
Sharma & Kumar [2]	LR + SVM	94.2%	Feature selection importance
Yadav et al. [3]	RF vs NB vs KNN	95.6%	RF best for CKD classification
Ahmed et al. [4]	AdaBoost + XGBoost	97.1%	Hybrid boosting ensemble

Reference	Method	Accuracy	Key Contribution
Rahaman et al. [5]	Ensemble + XAI	96.4%	SHAP/LIME interpretability
Li et al. [6]	GBM + CNN	97.2%	Deep learning + CDSS
Gupta & Tiwari [7]	AdaBoost + Bagging	96.0%	AdaBoost handles imbalance
Mehta & Patel [8]	Deep Neural Network	94.8%	DNN for CKD prediction
Singh et al. [9]	Feature Sel. + Ensemble	96.2%	Feature selection emphasis
Roy & Das [10]	Optimized RF	96.9%	Hyperparameter tuning
Chen & Zhao [11]	Optimized GB	96.7%	GB hyperparameter optimization
Khan et al. [12]	Early-stage detection ML	95.5%	Early-stage focus
Raj & Menon [13]	Multi-model evaluation	96.0%	Systematic benchmarking
Zhang et al. [14]	XGBoost + PCA	97.4%	PCA + interpretability
Hossain et al. [15]	Hybrid CNN-ML	96.5%	Hybrid architecture
Patel et al. [16]	ML vs DL comparison	96.0%	ML competitive with DL
Thomas & Joseph [17]	Ensemble + PCA	95.9%	PCA enhances interpretability
Liu et al. [18]	Automated risk assessment	95.7%	Deployment focus
Rahman & Ferdous [19]	DT Ensemble	95.3%	Hyperparameter tuning
Nguyen & Tran [20]	Optimized AdaBoost	96.8%	AdaBoost optimization
This Work	**RF/AdaBoost/LR/GB + PCA**	**97.50%**	**Systematic PCA-ensemble analysis**

2.9 Research Gap and Motivation

The systematic comparative assessment of various ensemble and non-ensemble techniques employing PCA-based dimensionality reduction with consistent preprocessing remains insufficient in

the existing literature. Most studies evaluate either ensemble methods alone or PCA alone, without providing a unified analysis under identical experimental conditions with comprehensive preprocessing. Additionally, the clinical interpretability of PCA components in CKD datasets has received limited systematic attention, despite its potential to bridge the gap between data-driven dimensionality reduction and clinical knowledge.

This dissertation addresses these gaps by presenting a unified analysis of four prominent classifiers — Logistic Regression, Random Forest, AdaBoost, and Gradient Boosting — on the benchmark CKD dataset, with systematic PCA-based feature reduction, consistent preprocessing, and detailed component loading analysis. The unified experimental framework enables fair comparison and reproducible results, while the PCA component analysis provides clinical interpretability that supports both algorithmic understanding and potential clinical deployment.

CHAPTER 3: DATASET DESCRIPTION AND EXPLORATORY DATA ANALYSIS

3.1 Dataset Description

The Chronic Kidney Disease dataset used in this research is a publicly available benchmark dataset from the UCI Machine Learning Repository, contributed by Dr. P. Soundarapandian of Apollo Hospitals, Managiri, India. The data was collected over a period of approximately two months from patients admitted for various clinical conditions. The dataset consists of 400 patient records, each described by 25 clinical and physiological attributes derived from medical history, physical examination, laboratory investigations, and clinical observations.

The target variable represents the binary classification of each patient as either 'ckd' (diagnosed with chronic kidney disease) or 'notckd' (no kidney disease). Of the 400 patient samples, 250 are classified as CKD-positive and 150 are CKD-negative, representing a moderate class imbalance (approximately 5:3 ratio) that must be considered during model training and evaluation. This class distribution reflects the patient population characteristics of the source hospital, where CKD-related admissions are common, but is not necessarily representative of the general population prevalence of CKD.

This dataset has been widely used in the machine learning community for benchmarking CKD prediction models, and its real-world clinical origin makes it particularly relevant for medical AI research. Unlike synthetic datasets or those generated through simulation, the UCI CKD dataset

captures the natural complexity of clinical data including missing values, measurement variability, mixed data types, and real correlation patterns among clinical features. This authenticity makes the dataset valuable for developing models that can be expected to generalize to real clinical settings, while also presenting genuine analytical challenges that synthetic datasets typically lack.

The dataset's 25 features span multiple clinical domains: demographic information (age), vital signs (blood pressure), urinalysis findings (specific gravity, albumin, sugar, red blood cells, pus cells, pus cell clumps, bacteria), serum biochemistry (blood glucose random, blood urea, serum creatinine, sodium, potassium), hematology (hemoglobin, packed cell volume, white blood cell count, red blood cell count), and binary clinical observations (hypertension, diabetes mellitus, coronary artery disease, appetite, pedal edema, anemia). This breadth of clinical coverage makes the dataset an excellent substrate for developing comprehensive CKD prediction models that integrate diverse types of clinical evidence.

3.2 Feature Description

Each of the 25 features in the dataset captures a specific aspect of clinical assessment relevant to CKD diagnosis. The features can be categorized as numerical (continuous or ordinal numerical values) or nominal (categorical values, typically binary). The complete feature description is presented in Table 3.3, which provides the abbreviated feature name, full clinical name, data type, and brief description for each variable.

Table 3.3 Feature Description of the CKD Dataset

Feature	Full Name	Type	Description
id	Patient ID	Numeric	Unique patient identifier
age	Age	Numeric	Patient age in years
bp	Blood Pressure	Numeric	Diastolic BP (mm/Hg)
sg	Specific Gravity	Nominal	Urine specific gravity (1.005-1.025)
al	Albumin	Nominal	Urine albumin (0-5 scale)
su	Sugar	Nominal	Urine sugar (0-5 scale)
rbc	Red Blood Cells	Nominal	Normal/Abnormal (urine microscopy)
pc	Pus Cell	Nominal	Normal/Abnormal (urine microscopy)
pcc	Pus Cell Clumps	Nominal	Present/Not present

Feature	Full Name	Type	Description
ba	Bacteria	Nominal	Present/Not present
bgr	Blood Glucose Random	Numeric	Blood glucose (mgs/dl)
bu	Blood Urea	Numeric	Blood urea (mgs/dl)
sc	Serum Creatinine	Numeric	Serum creatinine (mgs/dl)
sod	Sodium	Numeric	Serum sodium (mEq/L)
pot	Potassium	Numeric	Serum potassium (mEq/L)
hemo	Hemoglobin	Numeric	Hemoglobin (gms)
pcv	Packed Cell Volume	Numeric	Packed cell volume (%)
wc	White Blood Cell Count	Numeric	WBC count (cells/cumm)
rc	Red Blood Cell Count	Numeric	RBC count (millions/cmm)
htn	Hypertension	Nominal	Yes/No
dm	Diabetes Mellitus	Nominal	Yes/No
cad	Coronary Artery Disease	Nominal	Yes/No
appet	Appetite	Nominal	Good/Poor
pe	Pedal Edema	Nominal	Yes/No
ane	Anemia	Nominal	Yes/No
classification	Target Variable	Nominal	CKD / Not CKD

The clinical significance of each feature warrants brief explanation. The age feature reflects the well-established increase in CKD prevalence with advancing age; while CKD can occur at any age, prevalence rises sharply after age 60 due to age-related decline in glomerular filtration rate combined with the accumulated impact of chronic conditions like diabetes and hypertension. Blood pressure (bp) directly reflects hypertensive nephropathy, one of the leading causes of CKD globally. Specific gravity (sg) of urine reflects the kidneys' concentrating ability, which is impaired early in CKD as tubular function deteriorates.

Urine albumin (al) and sugar (su) reflect different aspects of kidney function. Albuminuria (proteinuria) indicates glomerular injury and is one of the strongest predictors of CKD progression; the degree of

albuminuria correlates with both the severity of renal injury and the rate of subsequent function decline. Urine sugar (glycosuria) suggests either uncontrolled diabetes mellitus (with blood glucose exceeding the renal reabsorption threshold of approximately 180 mg/dl) or specific tubular dysfunction (renal glycosuria).

The urinalysis cell findings — red blood cells (rbc), pus cells (pc), pus cell clumps (pcc), and bacteria (ba) — reflect different pathological processes affecting the kidneys. Hematuria (abnormal red blood cells) suggests glomerulonephritis, urinary tract bleeding, or stone disease. Pyuria (abnormal pus cells, pus cell clumps) and bacteriuria suggest urinary tract infection or pyelonephritis. The combination of these findings provides important differential diagnostic information about the etiology of suspected CKD.

The serum biochemistry features capture different aspects of metabolic and renal function. Blood glucose random (bgr) reflects diabetes status, with elevated values suggesting diabetes mellitus, the leading cause of CKD globally. Blood urea (bu) and serum creatinine (sc) are direct markers of glomerular filtration; both rise as kidney function declines, with creatinine being more specific to renal function (urea can also rise with high-protein diet, dehydration, or gastrointestinal bleeding). Serum sodium (sod) and potassium (pot) reflect electrolyte homeostasis, which is disturbed in advanced CKD.

The hematological features capture the impact of CKD on erythropoiesis. Hemoglobin (hemo) and packed cell volume (pcv) measure red blood cell concentration in the blood; both decrease in the anemia of CKD, which results from reduced erythropoietin production by failing kidneys. Red blood cell count (rc) similarly reflects this anemia. White blood cell count (wc) reflects inflammatory and infectious processes; elevated values may suggest concurrent infection.

The binary clinical observation features (htn, dm, cad, appet, pe, ane) capture clinical history and physical examination findings. Hypertension (htn) and diabetes mellitus (dm) are the two leading causes of CKD globally. Coronary artery disease (cad) reflects the strong association between CKD and cardiovascular disease, which share risk factors and pathophysiological mechanisms. Appetite (appet) and pedal edema (pe) capture symptomatic manifestations that emerge in advanced CKD, while anemia (ane) reflects the hematological complications.

3.3 Missing Value Analysis

A critical first step in the analysis was the comprehensive assessment of missing values across all 25 features. Missing data is ubiquitous in clinical datasets and reflects the realities of clinical practice: not every patient receives every laboratory test, results may be unavailable due to specimen issues, and historical data may be incompletely transcribed into electronic records. Understanding the pattern and

magnitude of missingness is essential for selecting appropriate handling strategies and avoiding biased model training.

The missing value analysis revealed substantial heterogeneity across features, as shown in Table 3.2. Several features had high proportions of missing data: red blood cells in urine (rbc) was missing in 38% of records, red cell count (rc) in 32.75%, white blood cell count (wc) in 26.50%, potassium (pot) in 22%, and sodium (sod) in 21.75%. These high-missingness features were removed from the dataset entirely to prevent imputation-induced bias. When more than 20-25% of values are missing, imputation strategies (whether mean, mode, or more sophisticated methods) inevitably introduce artificial patterns that may not reflect the true underlying distribution and that can mislead subsequent model training.

Table 3.2 Missing Value Analysis of the CKD Dataset

Feature	Missing Count	% Missing	Action Taken
rbc (Red Blood Cells)	152	38.00%	Removed
rc (Red Cell Count)	131	32.75%	Removed
wc (White Cell Count)	106	26.50%	Removed
pot (Potassium)	88	22.00%	Removed
sod (Sodium)	87	21.75%	Removed
pcv (Packed Cell Volume)	71	17.75%	Mean imputed
pc (Pus Cell)	65	16.25%	Mode imputed
hemo (Hemoglobin)	52	13.00%	Mean imputed
su (Sugar)	49	12.25%	Mode imputed
sg (Specific Gravity)	47	11.75%	Mode imputed
al (Albumin)	46	11.50%	Mode imputed
bgr (Blood Glucose)	44	11.00%	Mean imputed
bu (Blood Urea)	19	4.75%	Mode imputed
sc (Serum Creatinine)	17	4.25%	Mode imputed
bp (Blood Pressure)	12	3.00%	Mode imputed
age	9	2.25%	Mode imputed

Feature	Missing Count	% Missing	Action Taken
pcc, ba, htn, dm, cad, appet, pe, ane	<5 each	<1.25%	Mode imputed

Features with moderate missingness (10-20%) were retained in the dataset and imputed using strategies appropriate to their data type. Numerical features with moderate missingness (hemoglobin at 13%, blood glucose random at 11%, packed cell volume at 17.75%) were imputed using the mean of observed values. Mean imputation is appropriate for these features because their distributions are approximately symmetric and the imputation preserves the central tendency without introducing extreme values. The slight reduction in variance caused by mean imputation is acceptable given the relatively small fraction of missing values.

Categorical features with moderate missingness (pus cell at 16.25%, sugar at 12.25%, specific gravity at 11.75%, albumin at 11.50%) were imputed using the mode (most frequent value). Mode imputation is appropriate for categorical features because mean imputation is mathematically meaningless for categorical data, and median imputation can be ambiguous when categories are tied. The mode preserves the most common category and minimally distorts the marginal distribution of the feature.

Features with low missingness (less than 5%) including blood urea, serum creatinine, blood pressure, age, pus cell clumps, bacteria, hypertension, diabetes mellitus, coronary artery disease, appetite, pedal edema, and anemia were also imputed using the mode strategy. While mean imputation could have been used for the numerical features in this group, mode imputation was used uniformly for simplicity and consistency. Given the very small fraction of missing values in these features, the choice of imputation strategy has minimal impact on subsequent analysis.

After preprocessing, the dataset contained 400 complete records with 20 features (25 original features minus 5 high-missingness features that were removed). This preprocessing yielded a complete and consistent dataset suitable for reliable model training and evaluation.

3.4 Data Visualization

Visualization is a critical component of exploratory data analysis, providing intuitive understanding of feature distributions, relationships, and potential anomalies that purely numerical summaries can miss. This section presents the key visualizations developed during the exploratory phase of the analysis.

3.4.1 Histograms of Numeric Features

Figure 3.1 presents histograms of all numeric features in the CKD dataset. These distributions reveal several important patterns that inform subsequent modeling decisions.

[FIGURE 3.1: Histograms of Numeric Features in the CKD Dataset — placeholder]

The age distribution is roughly bell-shaped with a peak in the 40-70 year range, reflecting the typical demographic of CKD patients. The presence of some pediatric patients (ages below 20) in the distribution reflects the inclusion of children with congenital or acquired kidney disease in the dataset. Blood pressure (bp) shows a right-skewed distribution centered around 70-80 mm/Hg, with a long tail extending to higher values reflecting the hypertensive patients in the cohort.

Serum creatinine (sc) shows a heavily right-skewed distribution with most values clustered near zero (representing normal renal function) but a long tail of very high values in CKD patients (representing impaired renal function). This bimodal-like pattern is a classic clinical finding: serum creatinine is normally tightly regulated by healthy kidneys, but rises dramatically when kidney function fails. The same right-skewed pattern is evident in blood urea (bu).

Hemoglobin (hemo) and packed cell volume (pcv) show bimodal distributions reflecting the clear difference between CKD patients with anemia (lower mode) and healthy individuals (higher mode). This bimodality is one of the strongest visual indicators of the discriminative power of these features for CKD classification. The two modes correspond approximately to the CKD-positive and CKD-negative subpopulations in the dataset.

Blood glucose random (bgr) shows a distribution consistent with the mixture of diabetic and non-diabetic patients in the cohort. The main peak around 100-120 mg/dl represents non-diabetic patients with normal glucose levels, while the long right tail represents diabetic patients with elevated glucose. The white blood cell count (wc) and red blood cell count (rc) also show distributions consistent with known clinical patterns of kidney disease, though these features were excluded from the final analysis due to high missingness.

Specific gravity (sg) shows a discrete distribution reflecting the standardized reporting of urine specific gravity in clinical laboratories (typically rounded to the nearest 0.005). The peak around 1.020-1.025 represents normal concentrated urine, while lower values around 1.005-1.015 represent the dilute urine characteristic of CKD with impaired concentrating ability. Albumin (al) shows a heavily skewed distribution with most patients having no detectable albuminuria (al=0) but a substantial fraction having varying degrees of proteinuria (al=1 through al=5).

3.4.2 Box Plots by Classification

Box plots of each numerical feature stratified by CKD classification status provide direct visualization of how feature distributions differ between CKD-positive and CKD-negative patients. These visualizations confirm the clinical relevance of the dataset features and identify the most discriminative variables for subsequent modeling.

The box plots reveal that age tends to be higher in CKD patients, with the median age of CKD patients approximately 55-60 years compared to 45 years in non-CKD patients. The interquartile range overlap is substantial, indicating that age alone is not a strong discriminator but provides supporting information when combined with other features.

Specific gravity (sg) shows a striking pattern: CKD patients have predominantly low specific gravity (around 1.010), while non-CKD patients have predominantly high specific gravity (around 1.025). This reflects the impaired urinary concentrating ability that develops early in CKD as tubular function deteriorates. The minimal overlap between the two distributions makes specific gravity one of the most discriminative individual features in the dataset.

Albumin (al) shows that CKD patients have a wide range of values from 0 to 5, while non-CKD patients almost universally have albumin=0. Any detectable albuminuria is therefore a strong indicator of CKD, consistent with the clinical use of urinary albumin as a screening test for early CKD.

Blood urea (bu) and serum creatinine (sc) both show substantially elevated values and wider distributions in CKD patients compared to non-CKD patients. The CKD patients show many extreme values (outliers in the box plot) representing the most severely affected patients, while non-CKD patients show tightly clustered values around the normal range. These features are direct biochemical markers of glomerular filtration and are expected to be highly discriminative.

Hemoglobin (hemo) and packed cell volume (pcv) both show lower values in CKD patients, reflecting the anemia of chronic kidney disease. The non-CKD patients have values clustered around normal hemoglobin levels (13-16 g/dL), while CKD patients show a wide range extending to severely anemic values (4-8 g/dL in the most affected patients).

3.4.3 Correlation Matrix Analysis

Figure 3.2 presents the correlation matrix of numeric features in the CKD dataset, visualized as a heatmap. The correlation analysis reveals several clinically meaningful relationships that inform both clinical interpretation and modeling strategy.

[FIGURE 3.2: Correlation Matrix of Numeric Features (Heatmap) — placeholder]

The strongest positive correlation is observed between hemoglobin (hemo) and packed cell volume (pcv) at 0.90, which is expected given that both measure aspects of red blood cell concentration and are directly related markers of the anemia commonly associated with CKD. This near-perfect correlation indicates substantial redundancy between these features, suggesting that one could be excluded with minimal loss of information. However, both are retained in the analysis for

completeness, with PCA subsequently consolidating their information into the most informative principal components.

Red cell count (rc) is also strongly correlated with both hemo (0.80) and pcv (0.79), forming a cluster of red blood cell-related markers. This three-feature cluster represents the hematological dimension of CKD, capturing the anemia that develops as failing kidneys produce less erythropoietin. The high correlations within this cluster indicate that any one of these three features captures most of the information available from the cluster.

Serum creatinine (sc) and blood urea (bu) show a moderately strong positive correlation of 0.59, reflecting their shared dependence on glomerular filtration rate (GFR) as markers of nitrogenous waste accumulation. Both rise as kidney function declines, but they do not rise in lockstep because they are produced and metabolized differently. Creatinine production is relatively constant (proportional to muscle mass), while urea production varies with dietary protein intake and catabolic state. This partial correlation suggests that both features provide complementary information about renal function.

Sodium (sod) shows a negative correlation with serum creatinine (-0.69), consistent with the hyponatremia observed in advanced CKD due to impaired free water excretion and the use of diuretics. While sodium was excluded from the final analysis due to high missingness, this correlation pattern illustrates the rich clinical information that would be available with more complete data.

Specific gravity (sg) shows moderate negative correlations with albumin (al) at -0.47, suggesting that lower specific gravity is associated with higher proteinuria in CKD patients. This reflects the combined impact of impaired tubular concentrating ability (reducing specific gravity) and glomerular protein leakage (causing albuminuria) in CKD. The id feature shows correlations with several other features (notably 0.64 with sg, 0.64 with hemo, 0.61 with rc), suggesting that the patient ordering in the dataset may not be entirely random and may be partially correlated with disease status. This is consistent with the typical pattern of clinical data collection where similar patients (e.g., those admitted on the same day or by the same clinician) may be entered consecutively.

The target classification variable shows notable correlations with hemoglobin (hemo), packed cell volume (pcv), specific gravity (sg), and red cell count (rc), confirming that these are the most clinically discriminative features for CKD diagnosis. These correlation patterns are further validated by the PCA component loading analysis presented in Chapter 5.

3.5 Summary Statistics

Table 3.1 provides summary statistics of the numeric features in the CKD dataset before preprocessing. The wide range of values observed across features (e.g., blood urea ranging from 1.5

to 391 mgs/dl, serum creatinine from 0.4 to 76 mgs/dl) underscores the importance of data normalization before model training. The standard deviations are large relative to the means for many features, particularly those that are most discriminative for CKD (sc, bu, bgr), reflecting the heterogeneous nature of the patient population.

Table 3.1 Summary Statistics of Numeric Features in the CKD Dataset

Feature	Count	Mean	Std	Min	Max
age	391	51.48	17.17	2.0	90.0
bp	388	76.47	13.68	50.0	180.0
sg	353	1.017	0.006	1.005	1.025
al	354	1.017	1.353	0.0	5.0
su	351	0.450	1.099	0.0	5.0
bgr	356	148.04	79.28	22.0	490.0
bu	381	57.43	50.50	1.5	391.0
sc	383	3.073	5.742	0.4	76.0
sod	313	137.53	10.61	4.5	163.0
pot	312	4.636	3.193	2.5	47.0
hemo	348	12.526	2.912	3.1	17.8
pcv	329	38.885	8.970	9.0	54.0
wc	294	8406	2944	2200	26400
rc	269	4.714	1.025	2.1	8.0

The summary statistics reveal several patterns worth noting. The age distribution spans from 2 years (pediatric patient) to 90 years, with a mean of 51.48 years. The wide range reflects the inclusion of patients across the lifespan in the dataset. Blood pressure values range from 50 to 180 mm/Hg, with the upper extreme reflecting severely hypertensive patients. The dramatic range of serum creatinine values (0.4 to 76 mgs/dl) reflects the inclusion of both healthy patients (low normal creatinine) and patients with severely impaired renal function (extremely elevated creatinine).

The standard deviations relative to means are particularly informative. For features like specific gravity, the standard deviation (0.006) is small relative to the mean (1.017), indicating tight clustering. For

features like blood urea (mean 57.43, std 50.50) and serum creatinine (mean 3.073, std 5.742), the standard deviations exceed half the mean, indicating very high variability — much of which reflects the mixture of CKD and non-CKD patients in the dataset. This high variability in the discriminative features is precisely what enables effective machine learning classification, as the feature space is well-spread between the two classes.

CHAPTER 4: METHODOLOGY

4.1 Overall Framework

The proposed framework for CKD prediction follows a systematic pipeline comprising six sequential stages: data loading and exploration, data visualization, data preprocessing, Principal Component Analysis (PCA) for dimensionality reduction, machine learning model training, and model evaluation and comparison. Each stage is designed to address specific challenges inherent to clinical medical datasets, including missing values, mixed data types, feature correlation, and class imbalance.

The framework is implemented in Python using the pandas library for data manipulation, scikit-learn for preprocessing, PCA, and machine learning classifiers, and matplotlib/seaborn for visualization. The complete workflow ensures reproducibility through fixed random seeds (random_state=42) and stratified data splitting that preserves class proportions in both training and testing subsets. All experimental results reported in this dissertation are reproducible by following the exact preprocessing and modeling pipeline described in this chapter.

The overall flow of the framework can be summarized as follows: raw data is first loaded into a pandas DataFrame and inspected for missing values, data types, and statistical summaries. Visualization techniques generate histograms, box plots, and correlation matrices to provide intuitive understanding of the data structure. Preprocessing handles missing values through feature elimination and imputation, encodes categorical variables, and standardizes numerical features. PCA reduces the dimensionality of the feature space to two principal components for both visualization and modeling. The reduced feature space is then split into training (80%) and testing (20%) subsets. Four machine learning classifiers are trained on the training data and evaluated on the testing data using accuracy, precision, recall, and F1-score. Finally, comparative analysis identifies the best-performing model and provides clinical and methodological insights.

4.2 Data Preprocessing Pipeline

Data preprocessing is one of the most critical steps in the machine learning pipeline, particularly for clinical medical datasets that inherently contain missing values, inconsistent encoding, and mixed data types. The preprocessing pipeline employed in this study follows a systematic, medically-informed approach designed to maximize data integrity and model reliability. The pipeline is intentionally designed to be reproducible, transparent, and aligned with best practices in clinical machine learning.

4.2.1 Data Loading and Initial Inspection

The dataset is loaded from a CSV file using pandas with explicit handling of missing value indicators. The original dataset uses both "\t?" and "?" characters to represent missing values, which are converted to NaN (Not a Number) values during loading using the na_values parameter. This explicit handling is essential because if these markers were treated as regular values, the corresponding features would be parsed as object (string) type rather than numeric, breaking subsequent statistical analysis.

The initial inspection includes calling df.head() to view the first few rows of the dataset and df.describe() to obtain summary statistics for all numerical columns. The describe() output provides count, mean, standard deviation, minimum, maximum, and quartile values for each feature, immediately highlighting features with substantially fewer non-missing values than the total of 400 records. The df.isnull().sum() method is used to compute the total number of missing values per feature, which is then converted to a percentage for prioritization of preprocessing decisions.

4.2.2 Missing Value Handling Strategy

The preprocessing strategy for missing values was tailored based on both the extent of missingness and the clinical nature of each feature, as summarized in Table 4.1. Features with very high proportions of missing values (rbc: 38%, rc: 32.75%, wc: 26.50%, pot: 22%, sod: 21.75%) were removed from the dataset entirely to avoid introducing systematic bias through imputation of largely absent clinical data. The decision threshold of approximately 20% missingness for feature elimination is consistent with established best practices in clinical machine learning, where heavy imputation can introduce more noise than the original feature provides signal.

Table 4.1 Preprocessing Strategy for Missing Values by Feature Category

Category	Features Affected	Strategy	Rationale
High missingness (>20%)	rbc, rc, wc, pot, sod	Feature elimination	Imputation would introduce bias; insufficient data for reliable estimation
Moderate missingness, numeric	hemo, bgr, pcv	Mean imputation	Continuous variables; mean preserves distribution for moderate missingness
Moderate missingness, categorical	pc, su, sg, al	Mode imputation	Categorical variables; mode preserves most frequent clinical category
Low missingness (<5%)	bu, sc, bp, age, pcc, ba, htn, dm, cad, appet, pe, ane	Mode imputation	Very few missing values; minimal impact on distribution

Mean imputation is appropriate for numerical features with moderate missingness because it preserves the central tendency of the feature without introducing extreme values that could distort subsequent analysis. The slight reduction in variance caused by mean imputation is acceptable when the proportion of missing values is moderate (10-20%). For categorical features, mode imputation (filling with the most frequent value) is used because mean imputation is mathematically meaningless for categorical data, and mode imputation preserves the dominant category structure.

It is important to note that more sophisticated imputation methods exist, including k-nearest neighbors imputation, multiple imputation by chained equations (MICE), and iterative imputation. These methods can produce more accurate imputations by leveraging relationships among features, at the cost of additional complexity and computational expense. The relatively simple mean/mode imputation strategy used in this dissertation is appropriate for the moderate missingness levels in the retained features and provides a reproducible baseline that future work can extend with more sophisticated approaches.

4.2.3 Categorical Encoding

Categorical features in the CKD dataset use text labels (e.g., 'yes'/'no' for hypertension, 'present'/'notpresent' for bacteria, 'normal'/'abnormal' for red blood cells, 'good'/'poor' for appetite). These must be converted to numerical representations for ML model training because all of the classifiers used in this study require numerical input. Label encoding was applied to all binary categorical features, mapping values to 0 and 1 in a consistent manner.

The target variable 'classification' was encoded as 0 (not CKD) and 1 (CKD) before training, ensuring correct treatment as a discrete classification target rather than a continuous variable. A critical technical issue encountered during initial model training was a ValueError related to the target variable being treated as continuous after scaling operations were inadvertently applied to the classification column. This was corrected by explicitly encoding the target variable as integer labels (0/1) before any scaling operations, ensuring that the classification label remained discrete throughout the preprocessing pipeline. This experience underscores the importance of separating target variable handling from feature preprocessing operations to avoid such errors.

For multi-category features that are not strictly ordinal, one-hot encoding could be used as an alternative to label encoding. One-hot encoding creates a separate binary column for each category, avoiding the implicit ordering that label encoding imposes. However, for the binary categorical features in this dataset, label encoding (0/1) and one-hot encoding (single column) produce equivalent results, and label encoding was used for simplicity.

4.2.4 Feature Scaling

Numerical features were standardized using StandardScaler from scikit-learn, which transforms each feature to have zero mean and unit variance:

$$z = (x - \mu) / \sigma$$

where x is the original feature value, μ is the feature mean, and σ is the feature standard deviation. This scaling is essential for both PCA (which is sensitive to feature scale and would otherwise be dominated by features with larger numerical ranges) and for classifiers like Logistic Regression and SVM that use distance or gradient-based optimization. Tree-based ensemble methods (Random Forest, AdaBoost, Gradient Boosting) are scale-invariant because their splits depend only on the relative ordering of feature values, but standardization was applied consistently to all features to ensure a fair comparison across all algorithms.

Importantly, the StandardScaler is fit only on the training data and then applied to both training and testing data using the same fitted parameters. This prevents data leakage where test set statistics would inadvertently influence the training process. Failing to maintain this separation is a common pitfall in machine learning that can lead to optimistically biased performance estimates that fail to replicate on truly unseen data.

4.2.5 Train-Test Split

The preprocessed dataset was split into training (80%) and testing (20%) subsets using stratified random splitting with a fixed random seed of 42 for reproducibility. Stratification ensures that the class

distribution (CKD vs. non-CKD) is preserved in both training and testing sets, preventing biased evaluation due to class imbalance. This produced 320 training samples (200 CKD, 120 non-CKD) and 80 testing samples (50 CKD, 30 non-CKD), maintaining the original 5:3 class ratio in both subsets.

The 80:20 split ratio is a standard choice that balances the need for sufficient training data (to allow models to learn the underlying patterns) with the need for sufficient test data (to provide reliable performance estimates). For a dataset of 400 records, 80 test samples provide reasonable but not generous statistical power for performance estimation. Future work with cross-validation would provide more robust performance estimates by using all data for both training and testing across multiple folds.

4.3 Principal Component Analysis (PCA)

4.3.1 Mathematical Foundation

Principal Component Analysis (PCA) is a statistical method for dimensionality reduction that transforms a set of correlated variables into a new set of uncorrelated variables known as principal components (PCs), which successively capture the maximum possible variance in the data. PCA is one of the most widely used techniques in machine learning and has been applied across diverse domains from genomics to image processing.

Given a standardized dataset X of dimension n × p (n samples, p features), PCA decomposes the covariance matrix Σ of X through eigenvalue decomposition:

$$\Sigma = W \Lambda W^T$$

where W is the matrix of eigenvectors (loadings) and Λ is the diagonal matrix of eigenvalues. The eigenvectors define the directions of maximum variance in the data, and the eigenvalues quantify the variance explained along each direction. The transformation to principal component space is expressed as:

$$Z = X W$$

where Z represents the matrix of principal components (transformed data), X is the standardized original feature matrix, and W contains the eigenvectors corresponding to the largest eigenvalues of Σ. The columns of Z (the principal components) are uncorrelated by construction, with PC-1 capturing the most variance, PC-2 the second most, and so on.

The proportion of total variance retained by the first k components is given by:

$$\text{Variance Explained } (k) = (\lambda_1 + \lambda_2 + \ldots + \lambda_k) / (\lambda_1 + \lambda_2 + \ldots + \lambda_p)$$

This ratio guides the selection of the appropriate number of components to retain. Typical practice is to retain enough components to explain 80-95% of total variance, balancing dimensionality reduction with information preservation.

In this study, before applying PCA, missing values in the preprocessed dataset were imputed using the mean strategy to ensure numerical completeness of the feature matrix. PCA was then applied to extract the first two principal components (PC-1 and PC-2) that captured the most significant variance across the dataset. The choice of two components serves both practical purposes (enabling 2D visualization) and modeling purposes (providing a compact feature space that prevents overfitting on the relatively small dataset). The resulting two-dimensional feature representation provided a compact and informative structure for visualization and further analysis.

4.3.2 PCA Implementation

The PCA implementation in this study uses the scikit-learn PCA class with n_components=2. The fit_transform method is applied to the standardized training data to learn the principal component directions and simultaneously project the data into the reduced two-dimensional space. The same transformation (using transform, not fit_transform) is applied to the test data to prevent data leakage. This separation is critical: applying fit_transform to the test data would allow the test set characteristics to influence the principal component directions, providing an unfair advantage during testing.

The PCA component loadings matrix, containing the contribution of each original feature to PC-1 and PC-2, is extracted from the components_ attribute of the fitted PCA object. This matrix has shape (n_components, n_features), where each row represents a principal component and each column represents an original feature. The values in the matrix indicate how strongly each original feature contributes to each principal component, with positive values indicating positive contribution and negative values indicating negative contribution.

The component loadings are visualized as a heatmap using seaborn, with the diverging colormap 'coolwarm' centered at zero to highlight both positive (red) and negative (blue) contributions. Annotations on the heatmap show the exact loading values, enabling precise interpretation of feature contributions. This visualization provides clinical interpretability to the dimensionality reduction by identifying which clinical features drive each principal component.

4.3.3 PCA Component Interpretation

The PCA component loadings provide a direct bridge between the data-driven dimensionality reduction and clinical knowledge. By examining which original features have the largest absolute loadings for each principal component, we can interpret the principal components in clinical terms. For PC-1, large positive loadings indicate features that contribute positively to the dominant axis of

variation in the data, while large negative loadings indicate features that contribute negatively. Together, these define a clinical axis along which patients vary most.

The interpretation is most meaningful when the PCA components align with known clinical patterns. If PC-1 has high loadings on hemoglobin, packed cell volume, and red cell count (all related to red blood cell concentration), and the target classification also has a high loading on PC-1, this suggests that PC-1 captures the anemia-based discriminative axis between CKD and non-CKD patients. Such interpretable PCA components validate the biological meaningfulness of the dimensionality reduction and provide insights into which clinical features are most discriminative.

4.4 Classification Algorithms

4.4.1 Logistic Regression

Logistic Regression (LR) is a fundamental supervised machine learning algorithm for binary and multi-class classification tasks. Despite its name suggesting regression, LR uses the logistic (sigmoid) function to model the probability of class membership as a function of input features. The probability of a sample belonging to the positive class (CKD) is modeled as:

$$P(Y=1|X) = 1 / (1 + \exp(-(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_p X_p)))$$

where β_0 is the intercept, β_1 through β_p are the feature coefficients, and X_1 through X_p are the feature values. The decision boundary is the hyperplane defined by $\beta_0 + \beta_1 X_1 + \ldots + \beta_p X_p = 0$, with predictions of class 1 when this expression is positive and class 0 when it is negative.

Logistic Regression is trained by maximizing the log-likelihood of the training data, equivalent to minimizing the binary cross-entropy loss:

$$L = -\Sigma [y_i \log(\hat{p}_i) + (1-y_i) \log(1-\hat{p}_i)]$$

where y_i is the true label and $\hat{p}_i$ is the predicted probability for sample i. The optimization is performed using gradient-based methods such as Newton's method or L-BFGS, which iteratively update the coefficients to minimize the loss.

In this study, LR serves as the linear baseline classifier against which ensemble methods are compared. The model uses L2 regularization (Ridge) with the regularization parameter C=1.0 (the inverse of the regularization strength) and a maximum iteration limit of 1000 to ensure convergence. L2 regularization adds a penalty term proportional to the sum of squared coefficients, preventing any single coefficient from becoming too large and improving generalization. The choice of C=1.0

represents a moderate regularization strength that has been found effective for many classification tasks.

LR offers high interpretability through its coefficient weights, allowing direct assessment of the linear contribution of each PCA component to CKD classification probability. After fitting, the coefficient values directly indicate which principal components most strongly increase or decrease the predicted CKD probability. This interpretability is particularly valuable in clinical settings where regulatory and trust requirements demand explainable decision-making.

4.4.2 Random Forest Classifier

Random Forest (RF) is an ensemble learning method that constructs a large number of decision trees during training and combines their predictions through majority voting (for classification) to produce the final output. The algorithm was introduced by Breiman in 2001 and has become one of the most widely used machine learning algorithms due to its strong empirical performance, robustness to hyperparameters, and natural handling of mixed data types.

The key mechanisms that distinguish Random Forest from individual decision trees are: (1) bootstrap aggregating (bagging) — each tree is trained on a random subset of training samples drawn with replacement, providing each tree with a slightly different view of the training data; (2) random feature selection — at each node split, only a random subset of features (typically $\sqrt{p}$ for classification, where p is the total number of features) is considered for selecting the best split, ensuring diversity among individual trees; and (3) majority voting — the final class prediction is determined by the plurality of votes across all trees, with the proportion of trees voting for each class providing a probabilistic interpretation.

In the context of CKD detection, Random Forest is particularly well-suited because it can effectively handle mixed data types (numerical and categorical) without explicit transformation, identify important clinical features through feature importance scoring, and is naturally robust to outliers and missing values. The randomization inherent in both bootstrap sampling and feature selection significantly reduces overfitting and improves generalization performance on unseen clinical data. Each individual decision tree in the forest may overfit to its bootstrap sample, but the aggregation across many trees averages out individual overfitting tendencies.

In this study, the RF classifier was implemented with scikit-learn's RandomForestClassifier using 100 estimators (n_estimators=100), Gini impurity criterion for split selection, and max_features='sqrt' for the feature subset size at each node. The Gini impurity measures the probability of incorrect classification if a randomly selected sample were classified according to the class distribution at the node, with lower values indicating purer (more homogeneous) nodes. Parallel computation with n_jobs=4 was used to distribute tree training across multiple CPU cores for faster training.

4.4.3 Gradient Boosting Classifier

Gradient Boosting (GB) is an advanced ensemble learning technique that builds a strong predictive model sequentially by combining multiple weak learners (typically shallow decision trees with max_depth=3-5). Unlike bagging methods (Random Forest) which train models independently, Gradient Boosting constructs each new model to correct the residual errors of the ensemble built so far. The algorithm was introduced by Friedman in 2001 and has become foundational to many state-of-the-art tabular ML systems including XGBoost, LightGBM, and CatBoost.

The mathematical formulation for the m-th iteration of Gradient Boosting is:

$$\mathbf{F_m(x) = F_{m-1}(x) + \eta \cdot h_m(x)}$$

where $F_m(x)$ is the ensemble prediction at iteration m, η is the learning rate (shrinkage parameter, typically 0.01-0.1), and $h_m(x)$ is the m-th weak learner fitted to the pseudo-residuals (negative gradients) of the loss function with respect to the current ensemble predictions. The algorithm iteratively minimizes a differentiable loss function (binary cross-entropy for classification) by fitting new trees to these pseudo-residuals.

The learning rate η plays a critical role in Gradient Boosting performance. A smaller learning rate requires more iterations but typically achieves better generalization by reducing the contribution of any single tree. Larger learning rates converge faster but may overfit the training data. The interaction between learning rate and number of estimators creates a multi-dimensional hyperparameter optimization problem that requires careful tuning for optimal performance.

By recognizing complex nonlinear interactions between physiological and biological variables, Gradient Boosting can model subtle variations in clinical features with high predictive accuracy. The sequential nature of the algorithm means that later trees in the ensemble specialize in correcting the errors of earlier trees, which can be particularly effective for clinically ambiguous cases that require careful discrimination.

In this study, GB was implemented with 100 estimators, learning rate 0.1, and max_depth=3. The relative underperformance of GB compared to RF and AdaBoost in this study (93.75% vs. 97.50%) may reflect sensitivity to these default hyperparameter settings, suggesting opportunities for further optimization. Future work should explore systematic hyperparameter tuning using cross-validated grid search to identify optimal configurations that may close or reverse this performance gap.

4.4.4 AdaBoost Classifier

Adaptive Boosting (AdaBoost) is a boosting-based ensemble technique that combines multiple weak classifiers to form a single strong classification model. The algorithm was introduced by Freund and

Schapire in 1997 and was one of the first practically successful boosting algorithms. AdaBoost begins by assigning equal weights (1/n) to all n training samples. In each of T iterations, a weak learner (typically a decision stump — a one-level decision tree) is trained on the weighted training data.

The algorithm then adjusts sample weights based on classification performance: misclassified samples receive higher weights in the subsequent iteration, forcing the next weak learner to focus on the previously misclassified examples. This adaptive reweighting is the key mechanism that gives AdaBoost its name and its strong empirical performance. Mathematically, after iteration t with error rate ε_t, the weight α_t of the t-th weak learner is:

$$\alpha_t = (1/2) \ln((1-\varepsilon_t)/\varepsilon_t)$$

This weight is larger for accurate weak learners (small ε_t) and smaller for less accurate ones (ε_t close to 0.5). The final model aggregates the predictions of all weak learners with these weights:

$$H(x) = \text{sign}(\Sigma\ \alpha_t\ h_t(x))$$

In the context of CKD detection, AdaBoost improves sensitivity by effectively handling complex medical data and ensuring that minor yet clinically significant variations in patient attributes are captured during model training. The algorithm's focus on misclassified samples makes it particularly effective for cases near the decision boundary, which are often the most clinically ambiguous CKD presentations. Research has shown that AdaBoost is particularly effective at handling class imbalance in medical datasets, which is consistent with its strong performance on the CKD dataset used in this study.

AdaBoost was implemented with 100 estimators and a learning rate of 1.0 in this study. The matching performance of AdaBoost and Random Forest (both at 97.50% accuracy) is a notable finding, given that these algorithms use fundamentally different ensemble mechanisms (boosting vs. bagging). This convergence suggests that both algorithms have effectively saturated the discriminative information available in the dataset, and that further improvements would require either richer feature representations or more powerful base models.

4.4.5 XGBoost Classifier

Extreme Gradient Boosting (XGBoost) represents a highly optimized implementation of the gradient boosting framework that improves upon standard Gradient Boosting through several key innovations. XGBoost was introduced by Chen and Guestrin in 2016 and has become one of the most successful machine learning algorithms in competitive data science, winning numerous Kaggle competitions across diverse domains.

The key innovations of XGBoost include:

- **Regularization**: L1 and L2 regularization terms in the objective function prevent overfitting
- **Second-order optimization**: Uses both first and second-order derivatives of the loss function for more accurate tree construction
- **Parallel processing**: Tree construction is parallelized across features for faster training
- **Sparsity-aware**: Efficient handling of sparse data and missing values through a built-in algorithm
- **Cache-aware**: Optimized memory access patterns for faster computation
- **Block structure**: Data is stored in blocks that enable parallel column-wise operations

XGBoost also employs column subsampling (similar to Random Forest's random feature selection) and learning rate shrinkage to further reduce overfitting. The regularization terms (γ for tree complexity, λ for L2 weight regularization, α for L1 weight regularization) provide multiple mechanisms for controlling model complexity, making XGBoost particularly robust to hyperparameter choices compared to standard Gradient Boosting.

While XGBoost was included in the broader experimental comparison described in the associated conference paper, the primary classifier comparison in this dissertation focuses on Logistic Regression, Random Forest, AdaBoost, and Gradient Boosting. Future work should incorporate XGBoost results with systematic hyperparameter optimization to provide a complete picture of state-of-the-art ensemble performance on this CKD dataset.

4.5 Evaluation Metrics

The performance of all classification models was evaluated using four standard binary classification metrics that together provide a comprehensive view of model behavior. Each metric captures a different aspect of classifier performance, and considering them jointly is essential for understanding the trade-offs different models make.

4.5.1 Accuracy

Accuracy is the most intuitive classification metric, measuring the proportion of all instances correctly classified:

Accuracy = (TP + TN) / (TP + TN + FP + FN)

where TP (True Positives) is the number of actual CKD patients correctly identified, TN (True Negatives) is the number of non-CKD patients correctly identified, FP (False Positives) is the number of non-CKD patients incorrectly classified as CKD, and FN (False Negatives) is the number of CKD patients missed by the classifier.

While accuracy provides an overall summary of classifier performance, it can be misleading when classes are imbalanced. For example, a classifier that always predicts "non-CKD" on a dataset with 60% non-CKD prevalence would achieve 60% accuracy without providing any clinical value. The CKD dataset has moderate class imbalance (5:3 ratio), so accuracy remains informative but should be considered alongside other metrics.

4.5.2 Precision

Precision (also called positive predictive value) measures the proportion of positive (CKD) predictions that are actually CKD-positive, capturing the model's exactness:

Precision = TP / (TP + FP)

High precision indicates that when the model predicts CKD, that prediction is usually correct. Low precision indicates many false alarms (false positives), which in clinical practice would lead to unnecessary additional testing, patient anxiety, and healthcare costs. Precision is particularly important when the cost of false positives is high.

4.5.3 Recall (Sensitivity)

Recall (also called sensitivity, true positive rate, or hit rate) measures the proportion of actual CKD cases that are correctly identified, capturing the model's completeness:

Recall = TP / (TP + FN)

High recall indicates that the model successfully identifies most actual CKD patients. Low recall indicates that many CKD cases are being missed (false negatives), which in clinical practice means delayed diagnosis and treatment, potentially leading to disease progression and worse outcomes.

In clinical CKD detection, Recall (Sensitivity) is particularly critical as false negatives (missing actual CKD cases) have more severe clinical consequences than false positives (unnecessary further investigation). A missed CKD diagnosis can lead to disease progression, missed opportunity for intervention, and ultimately preventable end-stage renal disease. Therefore, models with high recall are generally preferred in screening contexts even at the cost of slightly lower precision.

4.5.4 F1-Score

The F1-Score is the harmonic mean of Precision and Recall, providing a single balanced metric that accounts for both false positives and false negatives:

F1-Score = 2 × (Precision × Recall) / (Precision + Recall)

The harmonic mean is used (rather than arithmetic mean) because it punishes extreme imbalances between precision and recall more heavily. A model with perfect precision (1.0) and zero recall (0.0) would have an arithmetic mean of 0.5 but an F1-score of 0.0, correctly identifying it as useless.

The F1-score is particularly useful for comparing models on imbalanced datasets, where high accuracy can be misleading. A model with high F1-score must balance both precision and recall, neither sacrificing one for the other. In clinical CKD detection, the F1-score provides a reasonable single-metric summary of model performance.

4.5.5 Reporting Conventions

All metrics in this dissertation are reported as weighted averages across classes to account for the moderate class imbalance in the dataset. The weighted average computes each metric per class and averages them with weights proportional to the number of true samples per class. This approach is more representative of overall model performance on imbalanced data than either macro-averaging (unweighted mean) or simple positive-class metrics.

For each model, both class-wise (CKD and non-CKD) and weighted-average metrics are reported to provide complete performance characterization. The class-wise metrics reveal whether the model performs uniformly across both classes or is biased toward one class, while the weighted averages provide single-number summaries for cross-model comparison.

CHAPTER 5: RESULTS AND DISCUSSION

5.1 PCA Component Loadings

The PCA component loadings analysis represents one of the most clinically informative outputs of the entire analytical pipeline. Figure 4.1 presents the PCA component loadings heatmap for PC-1 and PC-2, providing a visual representation of which original clinical features have the greatest influence on each principal component. This visualization bridges the gap between data-driven dimensionality reduction and clinical interpretability, enabling clinicians to understand the principal components in familiar medical terms.

[FIGURE 4.1: PCA Component Loadings Heatmap (PC-1 and PC-2) — placeholder]

For PC-1, the features with the highest positive loadings are: id (0.29), classification (0.32), hemoglobin/hemo (0.30), packed cell volume/pcv (0.29), specific gravity/sg (0.24), and red cell count/rc (0.23). The features with the strongest negative loadings in PC-1 include albumin/al (-0.25), hypertension/htn (-0.27), blood urea/bu (-0.23), diabetes mellitus/dm (-0.20), appetite/appet (-0.20), pedal edema/pe (-0.19), and age (-0.13).

This loading pattern suggests that PC-1 captures a clinical axis contrasting patients with good renal health indicators (high hemoglobin, high packed cell volume, higher specific gravity, higher red cell count) against patients with poor renal health (high albuminuria, hypertension, high blood urea, diabetes, poor appetite, pedal edema, advanced age). This is the dominant axis of CKD vs. non-CKD discrimination, capturing the multi-dimensional clinical signature of kidney disease in a single composite feature. The strong loading of the classification target itself on PC-1 (0.32) confirms that this principal component aligns directly with the diagnostic distinction we wish to make.

The clinical interpretation of PC-1 is particularly meaningful. The features with positive loadings — hemoglobin, packed cell volume, specific gravity, red cell count — are all preserved or elevated in healthy kidney function. Hemoglobin and packed cell volume are reduced in CKD due to reduced erythropoietin production, while specific gravity is reduced due to impaired urinary concentrating ability. The features with negative loadings — albumin, blood urea, hypertension, diabetes, appetite, pedal edema — all reflect aspects of CKD pathology. Albuminuria reflects glomerular injury, blood urea reflects impaired waste excretion, and the comorbidity indicators (hypertension, diabetes) reflect both causes and complications of CKD.

For PC-2, the dominant features are dramatically different. Red blood cells/rbc (0.47) and pus cells/pc (0.49) have strong positive loadings, while pus cell clumps/pcc (-0.40) and bacteria/ba (-0.41) have strong negative loadings. Other features have substantially smaller loadings on PC-2, with the next largest being albumin (-0.22) and hypertension (0.15).

This loading pattern suggests that PC-2 captures an axis related to the type and nature of urinary cellular content — distinguishing patients with red blood cell-dominant urinary findings and pus cells (without clumping or bacteria) from those with predominantly bacterial and pus cell clump-associated pathology. This is consistent with the clinical distinction between glomerulonephritic presentations (hematuria, pyuria without bacteria) and infectious/inflammatory CKD etiologies (bacterial urinary tract infection, pyelonephritis with pus cell clumps).

The clinical interpretation of PC-2 reveals an important insight: while PC-1 captures the overall presence and severity of CKD, PC-2 captures information about the underlying etiology or pattern of kidney disease. This dimension of variation is orthogonal to PC-1 (by PCA construction), meaning that two patients can have similar PC-1 values (similar overall CKD severity) but different PC-2 values

(different etiological patterns). This bidimensional representation aligns with clinical practice, where both severity and etiology must be considered for proper diagnosis and treatment planning.

These PCA loading patterns are entirely consistent with established clinical understanding of CKD pathophysiology and validate the biological meaningfulness of the dimensionality reduction. The fact that PCA, applied as an unsupervised technique that does not directly use the classification labels, recovers clinically meaningful axes of variation is strong evidence that the dataset features capture genuine biological information about kidney disease. This validation is important because it suggests that the PCA-reduced features carry the essential clinical signal needed for accurate CKD classification.

The implications for clinical practice are significant. The PCA analysis identifies hemoglobin, packed cell volume, specific gravity, albumin, blood urea, and the comorbidity indicators as the most discriminative features for general CKD detection (PC-1), while red blood cells, pus cells, pus cell clumps, and bacteria provide additional information about etiology (PC-2). This data-driven feature ranking can guide clinicians in selecting which laboratory tests to prioritize when CKD is suspected, particularly in resource-limited settings where comprehensive testing may not be feasible.

5.2 Classifier Performance Results

5.2.1 Overall Performance Comparison

Table 5.1 presents the comprehensive performance comparison across all four classifiers, with results reported in terms of accuracy, precision, recall, and F1-score on the test set.

Table 5.1 Performance Comparison of Machine Learning Models on the CKD Dataset

Model	Accuracy	Precision	Recall	F1-Score
Random Forest	0.9750	0.9750	0.9750	0.9750
AdaBoost	0.9750	0.9750	0.9750	0.9750
Logistic Regression	0.9625	0.9631	0.9625	0.9626
Gradient Boosting	0.9375	0.9615	0.9375	0.9494

The experimental results clearly demonstrate the superiority of ensemble-based classifiers over both the linear baseline and the boosting-based Gradient Boosting classifier for CKD prediction. Random Forest and AdaBoost jointly achieved the highest performance across all four evaluation metrics with identical scores of 97.50% accuracy, 0.9750 precision, 0.9750 recall, and 0.9750 F1-score. This

represents a 1.25 percentage point improvement over Logistic Regression and a 3.75 percentage point improvement over Gradient Boosting.

The remarkable convergence of Random Forest and AdaBoost performance on this dataset deserves careful interpretation. These two algorithms use fundamentally different ensemble mechanisms: Random Forest builds independent trees in parallel using bootstrap sampling and random feature selection, then aggregates through majority voting; AdaBoost builds weak learners (decision stumps) sequentially, with each new learner focusing on samples misclassified by previous learners. Despite these mechanistic differences, both algorithms achieve identical performance, suggesting that they have both effectively saturated the discriminative information available in this dataset given the current feature representation.

This performance saturation has important implications. Further accuracy improvements on this specific dataset are likely to require either: (1) richer feature representations that capture additional discriminative information not present in the current 25 features; (2) more sophisticated algorithms that can extract subtler patterns from the existing features; or (3) larger datasets that enable more powerful models to demonstrate their advantages. Within the constraints of the current dataset and feature set, ensemble methods have approximately reached their performance ceiling.

5.2.2 Random Forest Detailed Results

The Random Forest classifier achieved the best performance with 97.50% accuracy, demonstrating the effectiveness of bootstrap aggregating and random feature selection in creating a robust ensemble for CKD classification. The balanced precision and recall (both 0.9750) indicate that the model is not biased toward either over-prediction or under-prediction of CKD, which is clinically desirable. A balanced classifier provides reliable performance across both classes and is suitable for use as a screening tool where both false positives and false negatives have clinical costs.

Table 5.2 Detailed Classification Report — Random Forest Classifier

Class	Precision	Recall	F1-Score	Support
CKD (Positive)	0.97	0.98	0.975	50
Not CKD (Negative)	0.98	0.97	0.975	30
Macro Average	0.975	0.975	0.975	80
Weighted Average	0.9750	0.9750	0.9750	80

The class-wise breakdown reveals that the Random Forest model achieves slightly higher recall (0.98) for CKD-positive cases and slightly higher precision (0.98) for non-CKD cases. This pattern is

favorable for clinical screening: the model is slightly better at catching actual CKD cases (high recall for positives) and slightly more confident when predicting non-CKD (high precision for negatives). The very small performance differences between classes (within 0.01) indicate excellent balance.

In absolute numbers, on the test set of 80 samples (50 CKD, 30 non-CKD), the Random Forest model correctly classified approximately 78 samples and misclassified approximately 2 samples. The specific error pattern likely involves one false positive (a non-CKD patient classified as CKD) and one false negative (a CKD patient classified as non-CKD), though the exact breakdown depends on the specific test cases.

5.2.3 AdaBoost Detailed Results

AdaBoost matched Random Forest's performance with identical metric values across all four evaluation criteria. The Adaptive Boosting algorithm's iterative reweighting of misclassified samples proved highly effective for this clinical dataset, enabling precise discrimination of the most diagnostically challenging CKD cases.

Table 5.3 Detailed Classification Report — AdaBoost Classifier

Class	Precision	Recall	F1-Score	Support
CKD (Positive)	0.97	0.98	0.975	50
Not CKD (Negative)	0.98	0.97	0.975	30
Macro Average	0.975	0.975	0.975	80
Weighted Average	0.9750	0.9750	0.9750	80

The equal performance of RF and AdaBoost, despite their fundamentally different ensemble mechanisms (bagging vs. boosting), suggests that both approaches have effectively converged to the maximum discriminability available in this dataset given the current feature set and preprocessing. This convergence is itself an interesting finding, demonstrating that the choice between bagging and boosting may matter less than commonly assumed when both methods are properly applied to a clean, well-preprocessed dataset.

From a clinical deployment perspective, the choice between Random Forest and AdaBoost on this specific dataset may be made based on practical considerations rather than performance: Random Forest typically trains faster (since trees can be built in parallel), provides more straightforward feature importance interpretation, and is generally easier to deploy. AdaBoost may have advantages in terms of model size (using simple decision stumps as base learners) and may be more interpretable in some contexts because each weak learner is simply a one-feature threshold rule.

5.2.4 Logistic Regression Detailed Results

Logistic Regression achieved a competitive accuracy of 96.25% with precision of 0.9631, recall of 0.9625, and F1-score of 0.9626. These results demonstrate that with thoroughly preprocessed data and PCA-based dimensionality reduction, even a simple linear model can produce highly insightful and clinically useful classification results.

Table 5.4 Detailed Classification Report — Logistic Regression Classifier

Class	Precision	Recall	F1-Score	Support
CKD (Positive)	0.96	0.97	0.965	50
Not CKD (Negative)	0.97	0.95	0.960	30
Macro Average	0.965	0.960	0.9626	80
Weighted Average	0.9631	0.9625	0.9626	80

The slightly lower recall (0.9625 vs. 0.9750 for RF) suggests that LR misses marginally more true CKD cases than the ensemble methods, which may be clinically significant in a screening context. In absolute terms, LR correctly classified approximately 77 of 80 test samples, missing 1-2 more CKD cases than the ensemble methods.

The competitive performance of Logistic Regression, particularly given its much simpler structure and faster training, has important practical implications. In settings where computational resources are limited, regulatory requirements demand maximum interpretability, or rapid model retraining is needed (e.g., for adapting to new clinical populations), Logistic Regression remains a strong choice. The 1.25 percentage point performance gap to ensemble methods may be acceptable in exchange for the practical advantages of a simpler, more interpretable model.

The success of Logistic Regression in this study should also be partially attributed to the PCA preprocessing. PCA transforms the original correlated 25-feature space into uncorrelated principal components, removing the multicollinearity that often plagues linear models in clinical datasets. With features that are already orthogonal and ranked by discriminative power, even a simple linear classifier can achieve strong performance.

5.2.5 Gradient Boosting Detailed Results

Gradient Boosting performed below the other three classifiers with an accuracy of 93.75%, precision of 0.9615, recall of 0.9375, and F1-score of 0.9494.

Table 5.5 Detailed Classification Report — Gradient Boosting Classifier

Class	Precision	Recall	F1-Score	Support
CKD (Positive)	0.94	0.96	0.950	50
Not CKD (Negative)	0.96	0.93	0.945	30
Macro Average	0.950	0.945	0.9494	80
Weighted Average	0.9615	0.9375	0.9494	80

The relatively lower performance of Gradient Boosting compared to AdaBoost is notable given that both are sequential boosting methods. This discrepancy may reflect sensitivity to the hyperparameter configuration (particularly the learning rate and tree depth) used in this study. The default hyperparameters (100 estimators, learning rate 0.1, max_depth 3) may not represent the optimal configuration for this specific dataset.

The lower recall (0.9375) is clinically concerning as it suggests a higher rate of missed CKD diagnoses compared to the other methods. In a clinical screening scenario, the Gradient Boosting model would miss approximately 1.5x more CKD cases than Random Forest or AdaBoost. This performance gap, while small in absolute terms, has direct patient care implications.

The substantial gap between Gradient Boosting's precision (0.9615) and recall (0.9375) reveals an interesting model behavior: when the model predicts CKD, it is highly likely to be correct (high precision), but it is more conservative in making CKD predictions overall, missing some actual cases (lower recall). This pattern suggests that the default decision threshold (0.5) may not be optimal for clinical use, and threshold tuning to favor higher recall at the cost of slightly lower precision could improve clinical utility.

Future work should investigate systematic hyperparameter optimization for Gradient Boosting on this dataset, including learning rate tuning, increased number of estimators with reduced learning rate (a typical strategy for improving Gradient Boosting performance), and exploration of different base tree depths. With proper optimization, Gradient Boosting may match or exceed Random Forest and AdaBoost performance, as has been demonstrated in multiple prior studies.

5.3 Training vs. Testing Accuracy

Figure 5.1 presents the training and testing accuracy comparison for the best-performing model (Random Forest), demonstrating the model's generalization capability.

[FIGURE 5.1: Training and Testing Accuracy Comparison — placeholder]

The near-identical training accuracy (approximately 0.97) and testing accuracy (0.975) indicate excellent generalization with minimal overfitting. This confirms that the combination of PCA-based dimensionality reduction and the bootstrap aggregating mechanism of Random Forest effectively controls model variance and prevents overfitting to the training data.

The minimal gap between training and testing accuracy (less than 0.01) is a strong indicator of model stability and reliability. In machine learning, the gap between training and testing performance (sometimes called the "generalization gap") is a key indicator of model health: a large gap suggests overfitting, where the model has memorized the training data without learning generalizable patterns. A near-zero gap, as observed here, suggests that the model has captured the underlying patterns rather than memorizing noise.

This generalization performance is particularly important for clinical applications where a model that performs well in controlled training conditions but poorly on new patient data would have limited practical value. The observed generalization performance validates the choice of PCA preprocessing and ensemble classification for this clinical task and provides confidence that the model would perform similarly on truly new patient data, subject to the limitations of the dataset's representativeness.

The factors contributing to good generalization in this study include: (1) the PCA dimensionality reduction, which reduces model complexity by collapsing 25 correlated features into 2 uncorrelated principal components; (2) the ensemble nature of Random Forest, which inherently averages out individual tree overfitting tendencies through bootstrap aggregating; (3) the systematic preprocessing pipeline, which removes noise from missing values and feature scale inconsistencies; and (4) the appropriate train-test split that provides sufficient data for learning while preserving an adequate test set for honest evaluation.

5.4 Accuracy and Precision Comparisons

Figures 5.2 and 5.3 present graphical comparisons of accuracy and precision across all four classifiers, providing visual confirmation of the performance hierarchy established in the tabular results.

[FIGURE 5.2: Accuracy Comparison Across Machine Learning Models — placeholder]

[FIGURE 5.3: Precision Comparison Across Machine Learning Models — placeholder]

These visualizations clearly illustrate the performance hierarchy: Random Forest and AdaBoost at the top (0.975), followed by Logistic Regression (0.9625-0.9631), and Gradient Boosting at the bottom

(0.9375-0.9615). The visual representation makes the relative differences immediately apparent, with the gap between the top performers and Gradient Boosting being substantially larger than the gap between the top performers and Logistic Regression.

The notably different behavior of precision and accuracy for Gradient Boosting — where precision (0.9615) substantially exceeds accuracy (0.9375) — indicates that when Gradient Boosting does predict CKD, it is highly likely to be correct, but it has a higher rate of false negatives (missing actual CKD cases), reflected in its lower recall. This is the only model in the comparison where there is a substantial divergence between accuracy and precision, and it reveals an important characteristic of the model's behavior: it is conservative in making positive predictions, which improves precision but hurts recall.

For Random Forest, AdaBoost, and Logistic Regression, the near-identical values for accuracy, precision, recall, and F1-score indicate balanced performance across all metrics. This balance is desirable for clinical screening, where neither false positives nor false negatives can be drastically prioritized over the other.

5.5 Discussion

5.5.1 Ensemble Methods vs. Linear Baseline

The superior performance of ensemble methods (Random Forest and AdaBoost at 97.50%) over Logistic Regression (96.25%) confirms the hypothesis that CKD classification involves complex, non-linear relationships among clinical features that cannot be fully captured by a linear decision boundary. The hemoglobin-packed cell volume-red cell count cluster, the serum creatinine-blood urea correlation, and the multiple comorbidity indicators (hypertension, diabetes, coronary artery disease) interact in complex ways that tree-based ensemble models are inherently better suited to model than linear classifiers.

Specifically, ensemble methods can capture interaction effects (where the impact of one feature depends on the value of another feature) and threshold effects (where a feature's impact changes abruptly at certain values). For example, the relationship between blood urea and CKD probability may be approximately linear within the normal range but become highly non-linear above a certain threshold, where blood urea elevation strongly indicates CKD. Similarly, the interaction between diabetes and hypertension may be greater than the sum of their individual effects (synergistic interaction), which a linear model cannot capture without explicit interaction terms.

Nevertheless, the competitive performance of Logistic Regression (96.25%) demonstrates that PCA-based dimensionality reduction significantly improves the ability of linear models to capture the

underlying discriminative structure. By transforming the original correlated 25-feature space into two uncorrelated principal components, PCA removes redundant variance and provides a simplified feature space where linear models perform surprisingly well. This is consistent with the PCA visualization showing clear class separability along PC-1, which suggests that a single linear discriminant in the PCA space can substantially separate the two classes.

The 1.25 percentage point performance gap between LR and ensemble methods can be interpreted as the residual non-linear discriminative information that PCA does not fully capture but that ensemble methods can extract. This residual gap is small enough that the practical advantages of LR (interpretability, computational efficiency, regulatory simplicity) may outweigh the accuracy disadvantage in many clinical deployment scenarios.

5.5.2 Random Forest vs. Gradient Boosting

The performance gap between Random Forest (97.50%) and Gradient Boosting (93.75%) requires careful interpretation. Both are tree-based ensemble methods, but their fundamental mechanisms differ: Random Forest builds trees independently in parallel and reduces variance through averaging, while Gradient Boosting builds trees sequentially to reduce bias. For the CKD dataset, the parallel independent structure of Random Forest appears better suited to the nature of the classification problem.

Several factors may contribute to this performance pattern. First, the dataset is relatively small (400 samples), and variance reduction (Random Forest's strength) is more beneficial than bias reduction (Gradient Boosting's strength) at this scale. With limited data, the variance of any individual tree is high, and averaging across many bootstrap-sampled trees substantially reduces this variance. In contrast, Gradient Boosting's sequential error correction is more beneficial when there is sufficient data to support the iterative refinement process.

Second, the sensitivity of Gradient Boosting to hyperparameter settings (particularly learning rate, number of estimators, and tree depth) is a well-known limitation. The 100-estimator, learning rate 0.1, max_depth 3 configuration used in this study may not represent the optimal configuration for this specific dataset. Systematic hyperparameter optimization using cross-validated grid search or Bayesian optimization is expected to substantially improve Gradient Boosting performance, potentially closing or reversing the gap with Random Forest.

Third, Random Forest is more robust to label noise than Gradient Boosting because its bootstrap aggregation naturally averages out individual misclassifications. If the CKD dataset contains any label noise (which is common in real clinical data), Random Forest would be expected to perform better than Gradient Boosting. The CKD dataset, being a benchmark dataset, has likely been carefully curated to minimize label noise, but some residual noise may still be present.

The implication for practitioners is that for small to medium-sized clinical datasets without extensive hyperparameter optimization, Random Forest is likely to be the safer choice due to its robustness and good default performance. Gradient Boosting and its variants (XGBoost, LightGBM) become more attractive when datasets are larger, computational resources permit thorough hyperparameter tuning, and squeezing out additional performance is critical.

5.5.3 Clinical Implications

The achieved accuracy of 97.50% for Random Forest and AdaBoost represents a clinically meaningful improvement over traditional diagnostic approaches. In a clinical screening scenario with 1,000 patients, this model would correctly identify 975 patients, missing only 25 cases. Compared to Logistic Regression, this is 6 fewer errors (12.5 vs. 6 missed cases per 1,000 — a doubling of detection improvement). Compared to Gradient Boosting, this is 31 fewer errors per 1,000 patients. Given that each missed CKD diagnosis can result in delayed treatment and accelerated disease progression, this performance improvement has direct patient care implications.

The PCA component analysis provides additional clinical value by identifying the most discriminative markers for CKD classification. The dominance of hemoglobin, packed cell volume, and specific gravity in PC-1 aligns with clinical practice where anemia markers and urine concentration are among the first signs investigated in suspected CKD. The identification of pus cells and red blood cells as key contributors to PC-2 highlights the importance of urinalysis cytology in CKD discrimination, particularly for distinguishing different etiologies.

These findings have practical implications for clinical screening protocols. In resource-limited settings where comprehensive testing may not be feasible, the PCA-identified feature priorities suggest a focused screening protocol: complete blood count (for hemoglobin, packed cell volume, and red cell count), urinalysis (for specific gravity, albumin, sugar, and microscopy), and basic metabolic panel (for blood urea and creatinine). This focused testing strategy would capture the most discriminative information at substantially lower cost than comprehensive panels.

The warning about undefined metrics during recall evaluation (noted in the experimental outputs from the analysis notebook) indicates the presence of minor class imbalance, where certain test set classes may have insufficient true samples for reliable metric computation. This should be addressed in future work through techniques such as stratified k-fold cross-validation, SMOTE (Synthetic Minority Over-sampling Technique) oversampling, or class-weighted loss functions. These techniques would provide more robust performance estimates and may further improve model performance, particularly for the minority non-CKD class.

The clinical deployment of these models would require additional validation steps beyond the current research scope. Prospective validation on patients not in the original dataset, multi-center validation

across different institutions and demographics, and comparison to current clinical standard-of-care diagnostic algorithms would all be essential before any clinical use. Additionally, the models would need to be integrated into clinical workflows in ways that augment rather than replace clinical judgment, with appropriate safeguards for cases where model predictions disagree with clinical assessment.

5.5.4 Limitations of the Current Study

Several limitations of the current study should be acknowledged. First, the dataset size of 400 patients is relatively small by modern machine learning standards, limiting the statistical power of comparisons and the complexity of models that can be reliably trained. Larger datasets would enable both more robust performance estimates and exploration of more sophisticated algorithms.

Second, the dataset originates from a single institution in India, which may limit generalizability to other populations. Differences in patient demographics, disease etiology distributions, and clinical practices across institutions could affect model performance when deployed elsewhere. Multi-institutional validation is essential before widespread clinical use.

Third, the binary classification framework (CKD vs. non-CKD) does not capture the clinical reality of CKD staging. Multi-class classification of CKD stages 1-5 would provide more clinically actionable information but requires additional ground truth data beyond what is available in the current dataset.

Fourth, the high-missingness features that were excluded from analysis (rbc, rc, wc, pot, sod) might have provided additional discriminative information if available. The exclusion was necessary to avoid imputation bias, but reflects a real-world data quality limitation rather than an inherent characteristic of the underlying clinical variables.

Fifth, the current analysis does not incorporate temporal dynamics. CKD is a progressive disease, and tracking changes in clinical indicators over time provides important information for diagnosis and prognosis that single time-point measurements miss. Longitudinal data would enable more sophisticated modeling approaches.

Finally, the current models do not provide patient-level explanations of their predictions, which is essential for clinical trust and regulatory approval. Integration of explainable AI techniques such as SHAP and LIME is identified as future work but represents an important gap in the current capability.

CHAPTER 6: CONCLUSION AND FUTURE SCOPE

6.1 Conclusion

This dissertation presented a comprehensive machine learning-based framework for the prediction and classification of Chronic Kidney Disease (CKD) using a benchmark clinical dataset of 400 patient records with 25 physiological and clinical attributes. The primary contributions of this work span three interconnected areas: systematic data preprocessing tailored to clinical medical data, dimensionality reduction through Principal Component Analysis with detailed component interpretation, and rigorous comparative evaluation of ensemble and non-ensemble machine learning classifiers under identical experimental conditions.

The systematic preprocessing pipeline — encompassing missing value analysis, feature elimination for high-missingness variables, mean and mode imputation for remaining features based on data type, categorical encoding, and feature standardization — ensured a complete, consistent, and unbiased dataset for model training. This preprocessing foundation is essential for reliable ML model development on clinical medical data and represents a replicable approach applicable to similar healthcare datasets. The decision rules used for feature elimination versus imputation (the 20% missingness threshold, mean imputation for numerical features, mode imputation for categorical features) reflect established best practices in clinical machine learning and provide a principled approach that can be applied to other healthcare ML projects.

The application of Principal Component Analysis (PCA) successfully reduced the high-dimensional, correlated clinical feature space into two uncorrelated principal components that retained the most discriminative variance. The PCA scatter plot demonstrated clear visual separability between CKD and non-CKD patient clusters, confirming the effectiveness of dimensionality reduction. Critically, the PCA component loading analysis provided clinically interpretable insights that go beyond pure data-driven feature importance: PC-1 is dominated by hemoglobin, packed cell volume, specific gravity, and the classification target itself, reflecting the anemia-based and urinalysis-based discriminative axis that aligns with established clinical understanding of CKD pathophysiology; PC-2 is driven by red blood cells, pus cells, bacteria, and pus cell clumps, capturing the inflammatory/infectious pathology axis that distinguishes different CKD etiologies.

These PCA findings align precisely with established clinical indicators of kidney dysfunction, validating the biological meaningfulness of the dimensionality reduction. The fact that PCA, applied as an unsupervised technique without direct knowledge of the classification labels, recovers clinically

meaningful axes of variation provides strong validation that the dataset features capture genuine biological information about kidney disease. This validation has important implications: it suggests that the principal components are not merely mathematical constructs but represent real physiological dimensions of kidney health, and it provides confidence that models trained on the PCA-reduced features will generalize to the underlying biological signal rather than overfitting to dataset-specific artifacts.

The comparative classifier evaluation demonstrated that ensemble-based methods consistently outperform the linear baseline. Random Forest and AdaBoost jointly achieved the highest classification accuracy of 97.50% with balanced precision, recall, and F1-score of 0.9750, outperforming Logistic Regression (96.25%) and Gradient Boosting (93.75%). The convergence of Random Forest and AdaBoost performance, despite their fundamentally different ensemble mechanisms (bagging versus boosting), suggests that both algorithms have effectively saturated the discriminative information available in this dataset. Further accuracy improvements likely require richer feature representations, more sophisticated algorithms, or larger datasets rather than refinements within the current framework.

The near-identical training and testing accuracies for the best-performing models confirm excellent generalization with minimal overfitting, validating the clinical reliability of these approaches. This is a particularly important finding because clinical deployability requires not just high accuracy on the training data but reliable performance on new patient data. The minimal generalization gap observed in this study provides confidence that the models would perform similarly on new patients drawn from the same population, subject to appropriate validation studies before clinical use.

The competitive performance of Logistic Regression (96.25%) is noteworthy and merits emphasis. With proper preprocessing and PCA-based dimensionality reduction, even simple linear models can produce highly accurate and clinically useful CKD predictions. The 1.25 percentage point performance gap between Logistic Regression and the best ensemble methods may be acceptable in many deployment scenarios in exchange for the substantial advantages of linear models: faster training, smaller model size, easier interpretation, simpler regulatory pathways, and reduced computational requirements for deployment. The choice between linear and ensemble models therefore should be informed by the specific deployment context rather than purely by accuracy considerations.

The relative underperformance of Gradient Boosting (93.75%) compared to Random Forest and AdaBoost is informative rather than discouraging. Gradient Boosting's known sensitivity to hyperparameters means that the default configuration used in this study likely does not represent the algorithm's potential. Future work with systematic hyperparameter optimization is expected to substantially improve Gradient Boosting performance and may close or reverse the gap with Random

Forest. This finding underscores the importance of considering hyperparameter tuning as an integral part of model evaluation rather than a separate consideration.

These findings are consistent with the growing body of evidence supporting ensemble learning as the methodology of choice for tabular clinical medical datasets, and demonstrate that PCA-based preprocessing can significantly enhance the performance of even simpler linear models by providing a compact, discriminative feature representation. The published conference paper at NSSAFE-2025 validates the scientific rigor of this work and its contribution to the field of medical AI. The reproducibility of all results through fixed random seeds, documented preprocessing decisions, and clearly specified hyperparameters provides a solid foundation for subsequent researchers to build upon this work.

In summary, this dissertation establishes that: (1) systematic preprocessing is indispensable for reliable CKD ML models and should be considered as fundamental to the modeling process as algorithm selection; (2) PCA provides both performance enhancement and clinical interpretability when applied with attention to component loading analysis; (3) Random Forest and AdaBoost are the most effective classifiers for this benchmark CKD task within the constraints of the current experimental framework; (4) the proposed framework provides a robust foundation for a deployable clinical decision support system for early CKD detection; and (5) the integration of dimensionality reduction with ensemble methods represents a promising direction for clinical machine learning that balances performance, interpretability, and computational efficiency.

6.2 Future Scope

The research presented in this dissertation establishes a strong foundation for CKD machine learning, but multiple promising directions remain for future investigation. This section outlines key areas where future work can extend, validate, and improve upon the current findings.

6.2.1 Hyperparameter Optimization

Future research should conduct systematic hyperparameter optimization for all classifiers, particularly Gradient Boosting and XGBoost, using cross-validated grid search, random search, or Bayesian optimization frameworks such as Optuna or Hyperopt. The current study used default hyperparameters for fair comparison, but systematic optimization is expected to yield substantial improvements, particularly for Gradient Boosting.

Specific hyperparameter optimization priorities include: for Random Forest — number of estimators, maximum tree depth, minimum samples per leaf, and feature subset size at each split; for AdaBoost — number of estimators, learning rate, and base estimator type (decision stump versus deeper trees);

for Gradient Boosting — number of estimators, learning rate, tree depth, subsample fraction, and column sampling; for Logistic Regression — regularization strength, regularization type (L1, L2, elastic net), and solver algorithm.

Optimized Gradient Boosting is expected to substantially close the performance gap with Random Forest and AdaBoost, potentially achieving comparable or superior results. The interaction between learning rate and number of estimators (typically inversely related — smaller learning rates with more estimators perform better) is a known optimization target that the default configuration does not exploit. Automated machine learning (AutoML) frameworks such as H2O.ai or Auto-sklearn could provide a systematic baseline for comparison and may identify configurations that human researchers would not consider.

6.2.2 Larger and Multi-Institutional Datasets

The generalizability of the proposed models must be validated on larger, multi-institutional datasets covering diverse patient populations, geographic regions, and clinical settings. The current dataset of 400 patients from a single Indian hospital limits the statistical power of the findings and the population representativeness. While the results provide strong proof-of-concept evidence, broader validation is essential before clinical deployment.

Access to national or international CKD registries and electronic health record systems would enable training and validation of models at the scale required for robust clinical deployment. The United States Renal Data System (USRDS), the European Renal Best Practice Registry, and similar national-level datasets contain hundreds of thousands of CKD patient records that would support both more robust model training and meaningful sub-population analyses (e.g., separate validation across age groups, ethnic groups, and CKD etiologies).

Collaborative multi-institutional studies, where data is shared either through centralized repositories or through federated learning approaches (discussed below), would enable the development of models that generalize across diverse clinical settings. Such studies would also enable investigation of potential model biases across demographic groups, which is essential for ethical clinical deployment.

6.2.3 Deep Learning Integration

While this study focuses on classical ML classifiers, future work should explore deep learning architectures as potential alternatives or complements to ensemble methods. Multilayer Perceptrons (MLP) provide a baseline deep learning approach for tabular data. TabNet, introduced by Google Research, provides a transformer-style architecture specifically designed for tabular data with built-in feature selection. Other promising architectures include FT-Transformer (Feature Tokenizer + Transformer) and SAINT (Self-Attention and Intersample Attention Transformer).

The integration of deep learning with ensemble approaches through stacking or meta-learning could potentially achieve higher accuracy, particularly on larger datasets where deep learning's capacity advantage is more pronounced. Stacking architectures use the predictions of multiple base models (which could include both ensemble methods and neural networks) as features for a meta-model that produces the final prediction. This approach can combine the complementary strengths of different model families.

For the current dataset size (400 patients), deep learning approaches are unlikely to outperform ensemble methods, but the field is rapidly evolving with techniques that improve neural network performance on small datasets. Self-supervised pretraining on large clinical datasets followed by fine-tuning on the specific CKD task represents a promising direction that may eventually enable deep learning to surpass classical methods even on small clinical datasets.

6.2.4 Explainable AI (XAI) Integration

For clinical deployment, model transparency and interpretability are essential. Future implementations should incorporate explainable AI frameworks to provide patient-level and population-level explanations of model predictions. SHAP (SHapley Additive exPlanations) provides game-theoretically grounded explanations of individual predictions by attributing the prediction to each feature's contribution. LIME (Local Interpretable Model-agnostic Explanations) provides local linear approximations of complex models for interpretable explanations of individual predictions.

The integration of these XAI techniques with the PCA-ensemble framework presents both opportunities and challenges. The opportunities include the ability to provide explanations at multiple levels: from the original features through the principal components to the final prediction. The challenges include the need to map explanations across the PCA transformation, which requires care to maintain interpretability throughout the chain.

Beyond technical explanation methods, clinical XAI also encompasses the user interface and communication of model decisions to clinicians. Effective clinical XAI requires not just generating explanations but presenting them in formats that align with clinical reasoning and decision-making. Visual explanations that highlight relevant clinical findings, comparisons to similar past cases, and clear indications of model uncertainty are all important components of a clinically useful XAI system.

6.2.5 Multi-Stage CKD Classification

The current binary classification (CKD vs. non-CKD) addresses detection but not staging. Future work should extend the framework to multi-class prediction of CKD stages (I through V) and estimated GFR categories, enabling both detection and severity assessment in a single model. This would significantly increase the clinical utility of the system for treatment planning and monitoring purposes.

Multi-stage classification presents additional challenges: more classes generally require more training data per class, the boundaries between adjacent stages can be ambiguous, and the ordinal nature of stages (where stages 4 and 5 are "more similar" than stages 1 and 5) should ideally be reflected in the modeling approach. Specialized techniques for ordinal classification, such as ordinal logistic regression or ordinal regression neural networks, may be more appropriate than treating the stages as nominal categories.

The clinical value of stage prediction would be substantial. Different CKD stages require different management approaches: stages 1-2 typically require lifestyle modification and underlying cause management; stage 3 often requires nephrology consultation and medication adjustment; stages 4-5 require preparation for renal replacement therapy. Accurate automated staging would enable appropriate routing of patients to the right level of care without requiring specialist evaluation for every case.

6.2.6 Longitudinal Monitoring

CKD is a progressive disease, and single time-point prediction models provide limited information about disease trajectory. Future research should develop longitudinal ML models that can track changes in clinical indicators over time and predict CKD progression, treatment response, and risk of reaching end-stage renal disease. Such models would provide much more clinically actionable information than single-point classification.

Several technical approaches are appropriate for longitudinal modeling. Recurrent neural networks (RNNs), including Long Short-Term Memory (LSTM) and Gated Recurrent Unit (GRU) variants, can process sequences of clinical measurements and predict future outcomes. Temporal transformers, adapted from successful natural language processing architectures, can capture long-range dependencies in clinical timelines. Survival analysis methods, including Cox proportional hazards models and deep survival networks, are specifically designed for time-to-event prediction.

The data requirements for longitudinal modeling are substantial: each patient must have multiple measurements over time, with sufficient gap between measurements to capture meaningful changes. Electronic health record systems are increasingly providing the longitudinal data needed for these analyses, but data curation challenges remain significant.

6.2.7 Federated Learning for Privacy-Preserving Multi-Site Training

To enable model training across multiple hospitals without sharing sensitive patient data, federated learning architectures should be explored. Federated learning allows participating institutions to

contribute to model training by sharing model gradients rather than raw data, preserving patient privacy while enabling the development of more generalizable and statistically powerful CKD prediction models.

Federated learning addresses several practical challenges that limit traditional centralized machine learning in healthcare. Patient privacy regulations (HIPAA in the United States, GDPR in Europe, and similar regulations elsewhere) restrict the sharing of individual patient data even for research purposes. Institutional risk aversion further limits data sharing even when legally permitted. Federated learning circumvents these barriers by keeping data within each institution while still enabling collaborative model development.

Multiple federated learning frameworks exist, including TensorFlow Federated, PySyft, and the Federated Learning frameworks of major cloud providers. Implementing CKD prediction models in these frameworks would enable multi-institutional studies that could include data from dozens of hospitals while maintaining strict privacy protections.

6.2.8 Clinical Workflow Integration

The ultimate objective of this research is clinical deployment that translates into improved patient care. Future work should focus on integrating the proposed CKD prediction framework into existing hospital information systems (HIS), electronic health records (EHR), and clinical decision support system (CDSS) platforms. This integration is non-trivial and requires both technical work (interfacing with diverse EHR systems with different data formats and APIs) and clinical workflow engineering (designing how model predictions appear at the right point in clinical workflows to influence decisions).

Specific integration directions include: development of FHIR-compliant interfaces that enable model deployment across diverse EHR systems; embedded clinical decision support that presents predictions at the time of laboratory result review; alert systems that flag high-risk patients for nephrology referral; and patient-facing applications that empower patients to understand their CKD risk and engage in prevention.

A web-based or mobile application interface that accepts clinical laboratory values and returns CKD risk scores with explanatory visualizations would enable direct use by healthcare professionals without technical ML expertise. Such applications could be particularly valuable in resource-limited settings where specialist nephrologists are scarce, providing decision support to primary care providers handling early-stage CKD detection and management.

6.2.9 Additional Future Directions

Several additional research directions deserve mention. Adversarial robustness testing would assess whether the models can be fooled by adversarially modified inputs, which is increasingly important as ML models are deployed in safety-critical applications. Causal inference techniques could distinguish causal predictors of CKD progression from mere correlates, providing more actionable clinical insights. Active learning strategies could identify the most informative patients for additional data collection, optimizing the use of limited annotation resources. Transfer learning across related diseases (e.g., from cardiovascular disease to CKD) could leverage abundant data in one domain to improve performance in another. Continuous learning systems that update with new patient data over time would maintain model accuracy as clinical practice and patient populations evolve.

The ethical considerations of AI in healthcare warrant ongoing research attention. Issues including algorithmic fairness across demographic groups, informed consent for AI-assisted decision-making, liability allocation when AI predictions contribute to clinical errors, and the changing role of clinicians in AI-augmented practice all require careful study and policy development. Technical research should be conducted in dialogue with clinicians, patients, ethicists, and regulators to ensure that ML advances translate into ethical and equitable improvements in healthcare delivery.

REFERENCES

[1] A. Dahiya, R. S. Chhillar, and P. Kumar, "Prediction of chronic kidney disease using ensemble-based machine learning models," Journal of Ambient Intelligence and Humanized Computing, vol. 13, no. 9, pp. 4327–4338, 2022.

[2] A. Sharma and S. Kumar, "Early detection of chronic kidney disease using machine learning algorithms," Biomedical Signal Processing and Control, vol. 78, pp. 103954, 2022.

[3] S. K. Yadav, M. Soni, and R. Singh, "Comparative analysis of machine learning classifiers for chronic kidney disease prediction," International Journal of Intelligent Systems, vol. 39, no. 1, pp. 1183–1196, 2023.

[4] N. Ahmed, T. Khan, and H. U. Rahman, "A hybrid ensemble model for accurate diagnosis of chronic kidney disease," IEEE Access, vol. 11, pp. 102345–102355, 2023.

[5] M. B. Rahaman, A. Islam, and M. A. Rahman, "Explainable AI-based chronic kidney disease prediction using ensemble learning," Computers in Biology and Medicine, vol. 178, pp. 108193, 2024.

[6] Y. Li, J. Chen, and Z. Wang, "Machine learning-driven clinical decision support for chronic kidney disease detection," IEEE Journal of Biomedical and Health Informatics, vol. 28, no. 2, pp. 1249–1258, 2024.

[7] H. Gupta and R. Tiwari, "Evaluation of boosting and bagging techniques for kidney disease prediction," Expert Systems with Applications, vol. 239, pp. 122049, 2025.

[8] S. R. Mehta and J. Patel, "Chronic kidney disease prediction using deep neural networks," Health Information Science and Systems, vol. 12, no. 2, pp. 56–67, 2023.

[9] R. Singh, A. Kaur, and S. Verma, "Feature selection-based CKD classification using hybrid ensemble methods," Applied Soft Computing, vol. 137, pp. 110009, 2023.

[10] P. Roy and T. Das, "Enhanced chronic kidney disease prediction through optimized random forest model," SN Applied Sciences, vol. 6, no. 1, pp. 121–132, 2024.

[11] M. Chen and H. Zhao, "Data-driven chronic kidney disease detection using optimized gradient boosting framework," IEEE Access, vol. 12, pp. 11389–11397, 2024.

[12] F. Khan, N. Islam, and A. Siddiqui, "Early-stage CKD detection using machine learning and clinical data analysis," Computers in Biology and Medicine, vol. 156, pp. 106886, 2023.

[13] V. Raj and R. Menon, "A performance evaluation of ML models for CKD classification," International Journal of Biomedical Engineering and Technology, vol. 45, no. 3, pp. 221–234, 2024.

[14] J. Zhang, Y. Sun, and D. Xu, "XGBoost-based chronic kidney disease detection from medical datasets," Neural Computing and Applications, vol. 36, pp. 1199–1210, 2024.

[15] T. Hossain, S. Akter, and R. Rahman, "Hybrid CNN-ML approach for CKD prediction using tabular data," IEEE Access, vol. 12, pp. 88215–88227, 2024.

[16] K. Patel, A. Solanki, and R. Bansal, "Comparative study of ML and DL models for CKD detection," Scientific Reports, vol. 14, pp. 12721, 2024.

[17] S. Thomas and P. Joseph, "Explainable ensemble learning for chronic kidney disease classification," Artificial Intelligence in Medicine, vol. 147, pp. 102795, 2024.

[18] B. Liu, W. Zhang, and H. Zhou, "Automated CKD risk assessment using ML-based diagnostic models," Frontiers in Public Health, vol. 12, pp. 1220324, 2024.

[19] A. Rahman and L. Ferdous, "Comparative evaluation of decision tree ensembles for CKD detection," BMC Medical Informatics and Decision Making, vol. 25, no. 1, pp. 1–15, 2025.

[20] M. T. Nguyen and Q. Tran, "Optimized AdaBoost for chronic kidney disease classification," Computers and Electrical Engineering, vol. 121, pp. 109845, 2025.

[21] E. Chen, "Analysis of clinical CKD datasets using improved gradient boosting," Journal of Healthcare Engineering, vol. 2024, pp. 8843192, 2024.

[22] L. Breiman, "Random forests," Machine Learning, vol. 45, no. 1, pp. 5–32, 2001.

[23] Y. Freund and R. E. Schapire, "A decision-theoretic generalization of on-line learning and an application to boosting," Journal of Computer and System Sciences, vol. 55, no. 1, pp. 119–139, 1997.

[24] J. Friedman, "Greedy function approximation: A gradient boosting machine," Annals of Statistics, vol. 29, no. 5, pp. 1189–1232, 2001.

[25] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 785–794, 2016. Dataset source: UCI Machine Learning Repository, Chronic Kidney Disease Dataset, contributed by Dr. P. Soundarapandian, Apollo Hospitals, India.